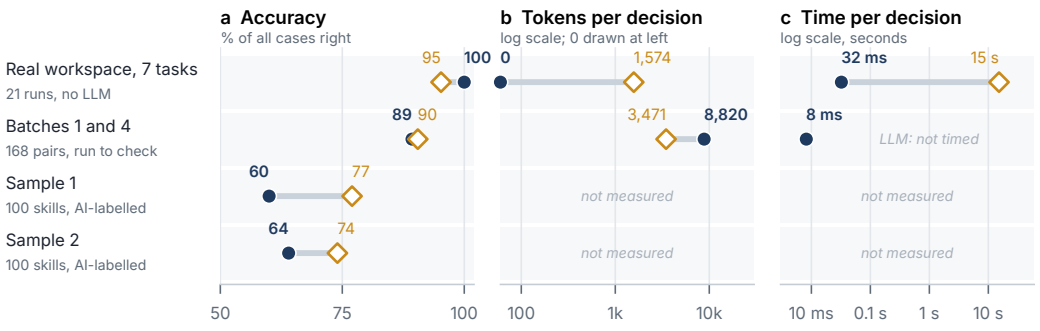
Real workspace, 7 tasks
21 runs, no LLM
Batches 1 and 4
168 pairs, run to check
Sample 1
100 skills, AI-labelled
Sample 2
100 skills, AI-labelled
a Accuracy
% of all cases right
95
100
89
90
60
77
64
74
50
75
100
b Tokens per decision
log scale; 0 drawn at left
0
1,574
3,471
8,820
not measured
not measured
100
1k
10k
c Time per decision
log scale, seconds
32 ms
15 s
8 ms
LLM: not timed
not measured
not measured
10 ms
0.1 s
1 s
10 s
skillc: verdict with a witness, or a named blocker
direct LLM verdict: an answer only