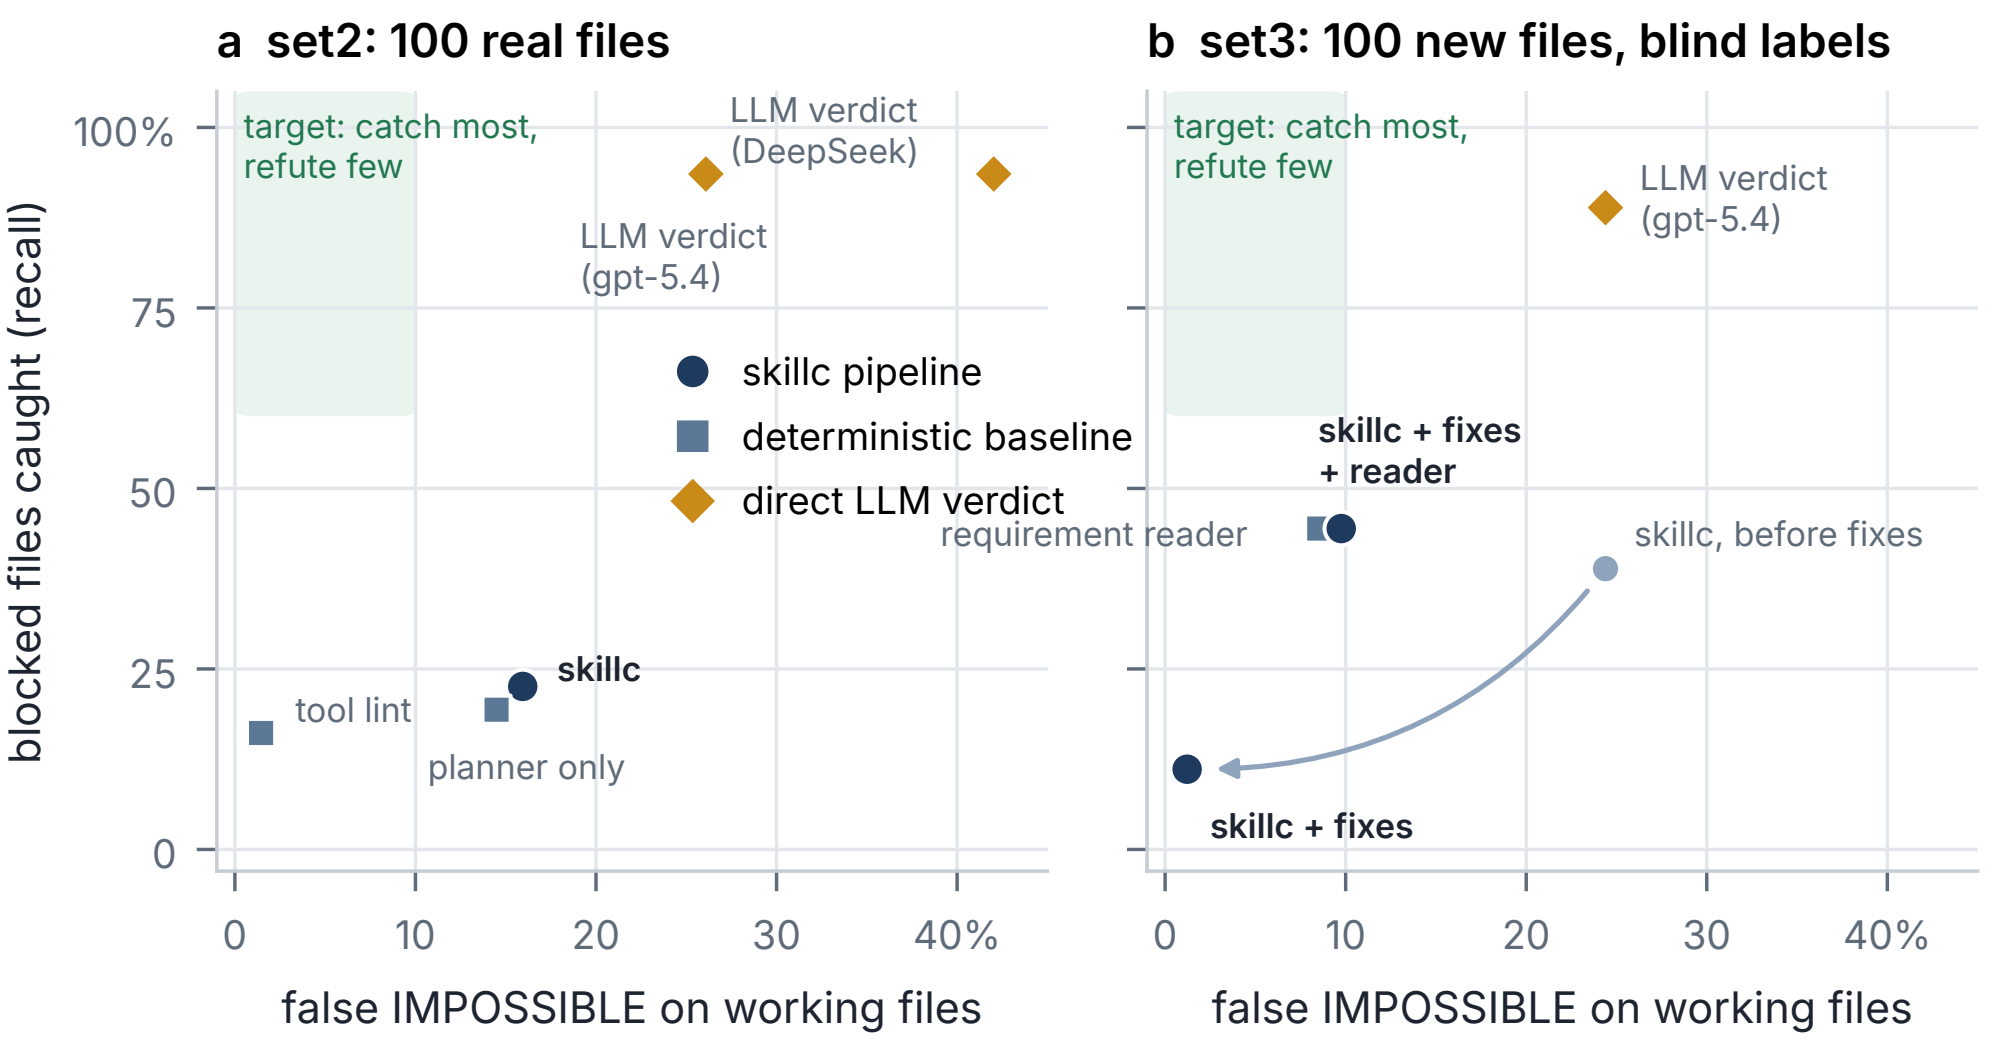
a set2: 100 real files
b set3: 100 new files, blind labels
blocked files caught (recall)
100%
75
50
25
0
target: catch most, refute few
LLM verdict (DeepSeek)
LLM verdict (gpt-5.4)
skillc pipeline
deterministic baseline
direct LLM verdict
skillc
tool lint
planner only
0
10
20
30
40%
false IMPOSSIBLE on working files
target: catch most, refute few
LLM verdict (gpt-5.4)
skillc + fixes + reader
requirement reader
skillc, before fixes
skillc + fixes
0
10
20
30
40%
false IMPOSSIBLE on working files