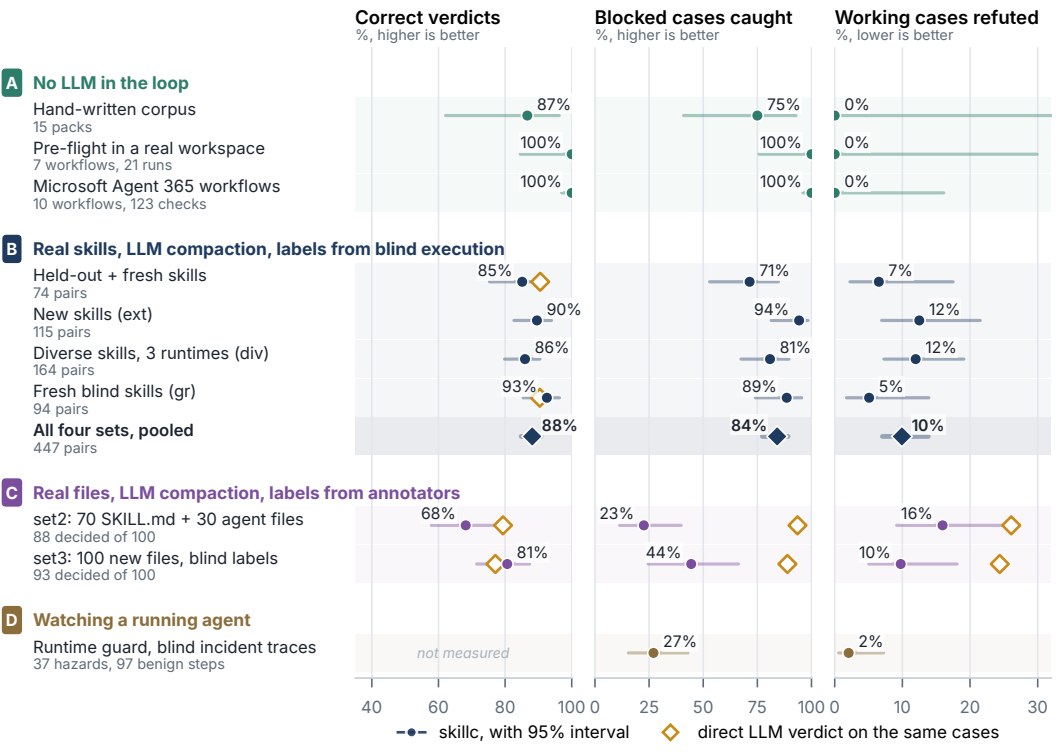
Correct verdicts
%, higher is better
Blocked cases caught
%, higher is better
Working cases refuted
%, lower is better
A
No LLM in the loop
Hand-written corpus
15 packs
87%
75%
0%
Pre-flight in a real workspace
7 workflows, 21 runs
100%
100%
0%
Microsoft Agent 365 workflows
10 workflows, 123 checks
100%
100%
0%
B
Real skills, LLM compaction, labels from blind execution
Held-out + fresh skills
74 pairs
85%
71%
7%
New skills (ext)
115 pairs
90%
94%
12%
Diverse skills, 3 runtimes (div)
164 pairs
86%
81%
12%
Fresh blind skills (gr)
94 pairs
93%
89%
5%
All four sets, pooled
447 pairs
88%
84%
10%
C
Real files, LLM compaction, labels from annotators
set2: 70 SKILL.md + 30 agent files
88 decided of 100
68%
23%
16%
set3: 100 new files, blind labels
93 decided of 100
81%
44%
10%
D
Watching a running agent
Runtime guard, blind incident traces
37 hazards, 97 benign steps
not measured
27%
2%
40
60
80
100
0
25
50
75
100
0
10
20
30
skillc, with 95% interval
direct LLM verdict on the same cases