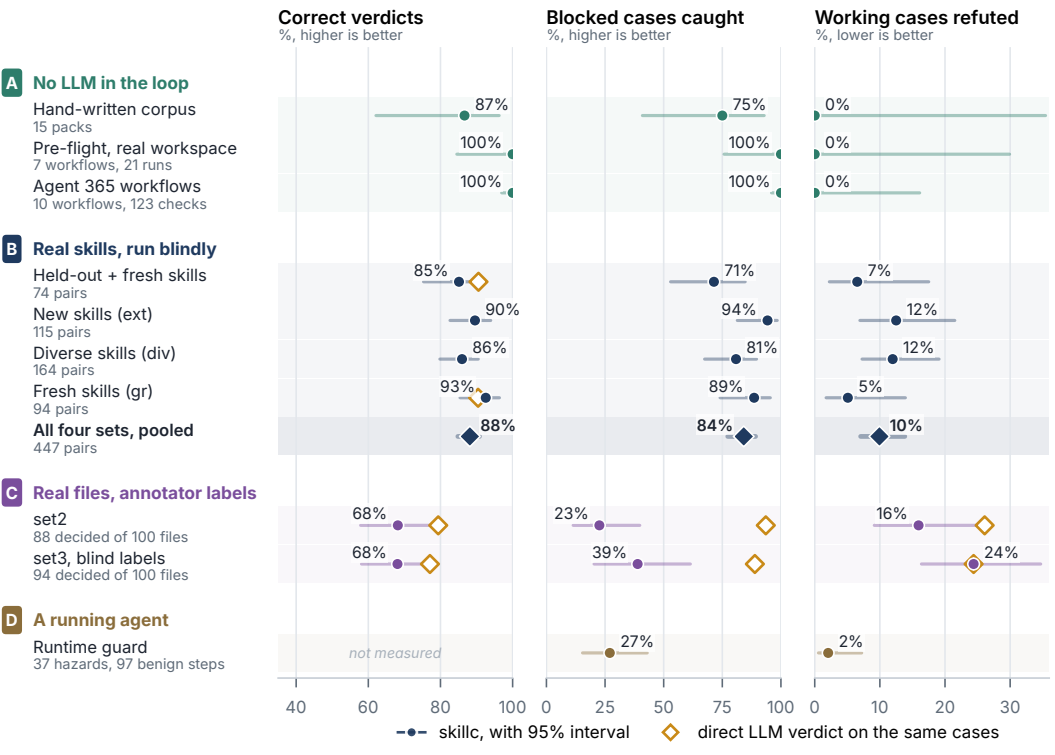
Correct verdicts
%, higher is better
Blocked cases caught
%, higher is better
Working cases refuted
%, lower is better
A No LLM in the loop
Hand-written corpus
15 packs
87%
75%
0%
Pre-flight, real workspace
7 workflows, 21 runs
100%
100%
0%
Agent 365 workflows
10 workflows, 123 checks
100%
100%
0%
B Real skills, run blindly
Held-out + fresh skills
74 pairs
85%
71%
7%
New skills (ext)
115 pairs
90%
94%
12%
Diverse skills (div)
164 pairs
86%
81%
12%
Fresh skills (gr)
94 pairs
93%
89%
5%
All four sets, pooled
447 pairs
88%
84%
10%
C Real files, annotator labels
set2
88 decided of 100 files
68%
23%
16%
set3, blind labels
94 decided of 100 files
68%
39%
24%
D A running agent
Runtime guard
37 hazards, 97 benign steps
not measured
27%
2%
40
60
80
100
0
25
50
75
100
0
10
20
30
skillc, with 95% interval
direct LLM verdict on the same cases