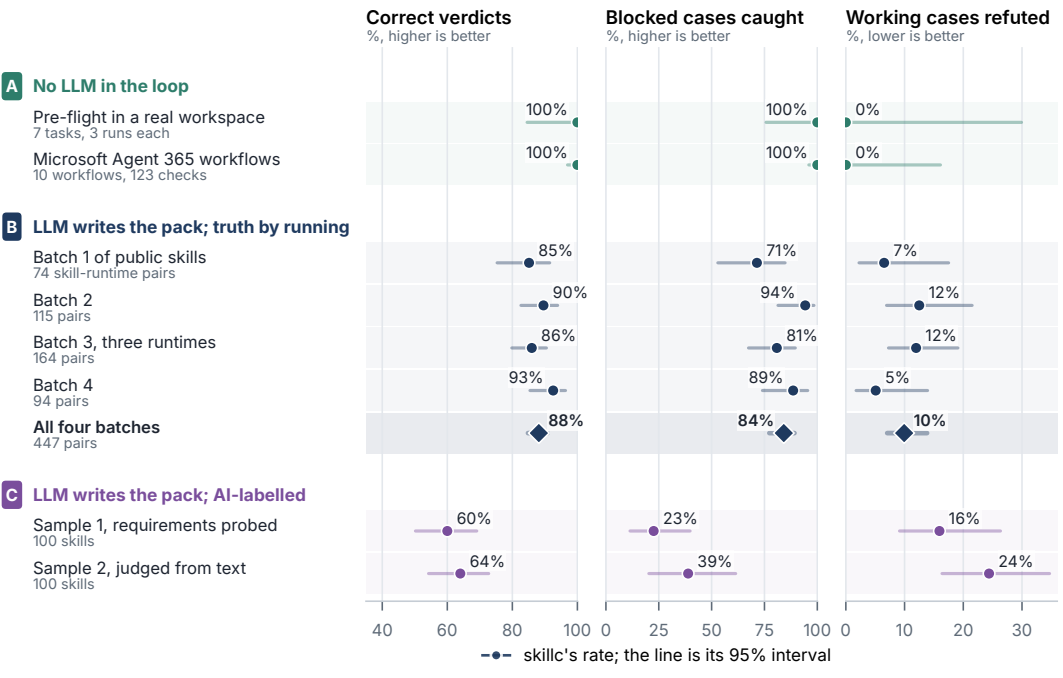
Correct verdicts
%, higher is better
Blocked cases caught
%, higher is better
Working cases refuted
%, lower is better
A
No LLM in the loop
Pre-flight in a real workspace
7 tasks, 3 runs each
100%
100%
Microsoft Agent 365 workflows
10 workflows, 123 checks
100%
100%
0%
0%
B
LLM writes the pack; truth by running
Batch 1 of public skills
74 skill-runtime pairs
85%
71%
7%
Batch 2
115 pairs
90%
94%
12%
Batch 3, three runtimes
164 pairs
86%
81%
12%
Batch 4
94 pairs
93%
89%
5%
All four batches
447 pairs
88%
84%
10%
C
LLM writes the pack; AI-labelled
Sample 1, requirements probed
100 skills
60%
23%
16%
Sample 2, judged from text
100 skills
64%
39%
24%
40
60
80
100
0
25
50
75
100
0
10
20
30
skillc's rate; the line is its 95% interval