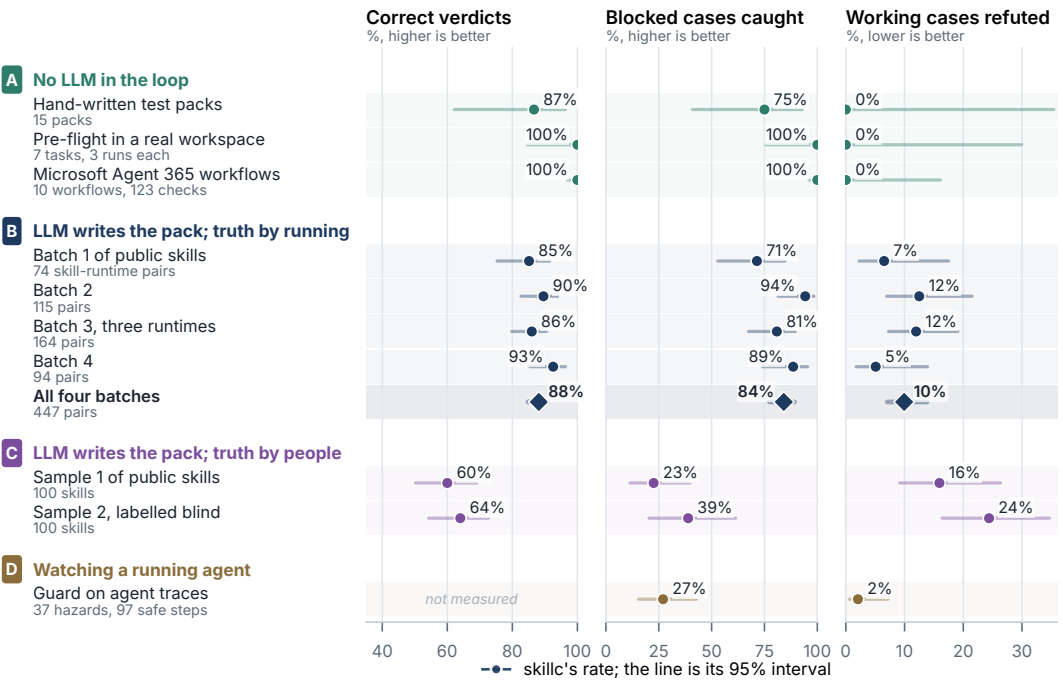
Correct verdicts
%, higher is better
Blocked cases caught
%, higher is better
Working cases refuted
%, lower is better
A
No LLM in the loop
Hand-written test packs
15 packs
Pre-flight in a real workspace
7 tasks, 3 runs each
Microsoft Agent 365 workflows
10 workflows, 123 checks
B
LLM writes the pack; truth by running
Batch 1 of public skills
74 skill-runtime pairs
Batch 2
115 pairs
Batch 3, three runtimes
164 pairs
Batch 4
94 pairs
All four batches
447 pairs
C
LLM writes the pack; truth by people
Sample 1 of public skills
100 skills
Sample 2, labelled blind
100 skills
D
Watching a running agent
Guard on agent traces
37 hazards, 97 safe steps
87%
100%
100%
75%
100%
100%
0%
0%
0%
85%
90%
86%
93%
88%
71%
94%
81%
89%
84%
7%
12%
12%
5%
10%
60%
64%
23%
39%
16%
24%
not measured
27%
2%
40
60
80
100
0
25
50
75
100
0
10
20
30
skillc's rate; the line is its 95% interval