
Can This Agent Even Do That?

A Decidable Goal-Achievability Type Discipline for LLM-Synthesized Agent Skills

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 LLM agents are increasingly driven by natural-language specifications, expressed
2 through artifacts such as `SKILL.md` and `AGENT.md` files that describe an agent’s
3 capabilities, tools, and workflows. Agent Skills, for example, package instructions
4 that guide an agent in performing specific tasks [5]. But how can we know whether
5 such specifications are themselves correct? Even if an LLM agent follows a speci-
6 fication exactly, can it actually achieve the goal that the specification is intended to
7 accomplish? Or might the goal be fundamentally unachievable, discovered only
8 after spending substantial time and tokens executing the agent—or even ending
9 in deadlock? We present a *skill compiler* that formally validates, before runtime
10 execution and without relying on an LLM judge, whether a skill’s *goal is achiev-*
11 *able*. The compiler combines capability-guarded reachability from automated
12 planning with deadlock-freedom from **simple multiparty sessions (SMPS)**, decided
13 *directly* over session configurations and a goal-marked global type synthesized
14 from the natural-language specification. Our approach follows the **SMPS** discipline
15 of checking a session directly against a global protocol [1, 2]. The approach begins
16 by using an LLM to extract the content of a natural-language specification and
17 compile it into a formal logical pack consisting of preconditions, effects, and goals,
18 which are then consumed by the inference rules. The checker is specified by a
19 Coq-verified, refutation-sound abstraction theorem. We further establish subject
20 reduction and session fidelity for the direct discipline. The single-label commu-
21 nication fragment with bystander interleaving is mechanized, while world-changing
22 action interleaving is proved under explicit hypotheses of goal stability and effect
23 commutativity. We formally define the language and its behavior, specify how
24 skills and sessions are checked, and prove that the resulting checking framework
25 has several important correctness properties. We also identify which part of the
26 system remains decidable and show where the boundary of undecidability begins
27 when participants can be added dynamically.

28 1 Introduction

29 A modern agent is no longer merely a chatbot, but an autonomous system capable of taking actions
30 beyond text generation—for example, writing code, sending emails, or executing trades. Such
31 behaviour is often guided by natural-language artifacts that specify instructions, workflows, and
32 constraints for the agent [3, 4, 6, 7]. A *skill* is one common way of expressing such guidance. It is a
33 folder containing instructions and tooling that an agent can discover and load as needed [5], thereby
34 *harnessing* a general-purpose agent toward a particular outcome.

35 There is growing evidence that this guidance matters. Across seven state-of-the-art open-source
36 multi-agent systems, failure rates range from 41% to 86.7% [8]. More importantly, some of these
37 failures can be addressed without changing the underlying model. In one system, simply adding a
38 task-objective verification step to the specification improves task success by 15.6% [8, Appendix H].
39 Likewise, repairing the harness based on execution traces improves task completion by 6.3 to 18.4
40 percentage points across four agent benchmarks [4]. These results show that the instructions and
41 structure surrounding an agent can matter as much as the model itself.

42 The question, however, is how we can know whether such guidance is sufficient to achieve the
43 outcome it declares. A plan may *book a flight and email a confirmation* even though no email tool is
44 available. It may pursue *a flight under \$500* even though the available booking tool can return only
45 premium fares. Or a planner may wait for a worker’s result while the worker waits for the planner’s
46 answer, with each expecting a message that the other will never send. These are not merely execution
47 failures. In each case, the goal is already unreachable before execution begins.

48 Such failures are structural and recurring, yet today the only way to discover them is often to run
49 the skill and inspect what went wrong afterwards [4]. By then, the agent may already have spent
50 substantial time and tokens taking actions, and a multi-agent system may already have become stuck
51 waiting for communication that will never occur. A simple static check could, in principle, identify
52 these problems before execution.

53 This paper asks a deliberately narrow question and answers it with a deterministic formal typing
54 system: **given an LLM skill, can its goal be achieved at all?** We do not, here, verify that every step
55 is safe, nor that the agent behaves correctly on every possible input. Instead, we decide *achievability*.
56 The analysis is tolerant: small details, detours, and payload variation should not trigger a false alarm.
57 At the same time, it is *sound for refutation*: when the compiler concludes that a goal is impossible,
58 that verdict is correct.

59 1.1 The idea, and the trust boundary

60 The architecture, shown in Figure 1, has two parts separated by a trust boundary. First, an LLM
61 performs *compaction*: it reads the natural-language skill and turns it into a formal *pack* containing
62 capabilities with preconditions and effects, a goal-marked global protocol, and an initial state. This
63 step is *untrusted*: the LLM may misunderstand the skill or omit relevant details.

64 A deterministic *checker* then decides achievability using only the resulting pack. The checker is
65 mechanically proved sound: if it returns *impossible*, then the goal is impossible under the formal pack
66 it received. If the LLM omits a constraint or capability from the original skill, the checker cannot
67 recover information that is not present in the pack; it will instead reason about the incomplete pack.
68 Conversely, if the LLM introduces a constraint that was not in the original skill, the checker may
69 report the resulting goal as impossible. Thus, our guarantee is about the checker’s verdict with respect
70 to the formal pack, not about whether the pack perfectly represents the original natural-language skill.

71 This does not make the compaction step a weakness that we simply ignore. The formal pack provides
72 a structured target for the LLM to extract. Rather than asking the LLM to freely invent a formal
73 model, we give it a predefined logical structure: capabilities have preconditions and effects, goals are
74 explicit, and the protocol follows a fixed syntax. This makes compaction a constrained translation
75 task and gives us a concrete artifact that can be inspected before execution.

76 Crucially, the gap between *what the user meant* and *what was formalized* cannot be eliminated. We
77 instead make this gap explicit at the compaction step, where the resulting pack can be inspected and
78 checked before execution. When the compaction is accurate, the checker can identify goals that are
79 impossible before the agent spends time and tokens executing the skill. When the compaction is
80 inaccurate, the checker still provides a sound verdict for the formal pack it received. The key design
81 idea is therefore to move achievability checking from runtime execution to an explicit, inspectable
82 checkpoint before execution.

83 1.2 Contributions

- 84 1. **A goal-achievability type discipline** (Section 3 and Appendices A–C) that combines
85 capability-guarded reachability from automated planning with deadlock-freedom from
86 multiparty session types. The discipline checks whole session configurations directly against




Figure 1: The trust boundary. An untrusted LLM compacts natural-language instructions into a formal pack. A deterministic checker then decides achievability over the pack using a Coq-verified abstraction theorem. The verdicts are asymmetric: *impossible* is sound (Theorem 1), while *achievable* depends on the accuracy of the pack (Section 3.1).

- 87 a goal-marked global type synthesized from the skill. Building on the multiparty session
 88 type tradition of Honda, Yoshida, and Carbone [12, 13], and the direct session-checking
 89 approach of **SMPS** [1, 2], it adds capabilities, explicit goals, and a refutation-sound but
 90 achievement-incomplete analysis for LLM-drafted skills.
- 91 2. **A mechanized soundness specification** (Appendix D) formalized in Coq. We prove a
 92 refutation-sound abstraction theorem parameterized by simulation hypotheses, together
 93 with tolerance soundness and capability monotonicity: adding capabilities cannot make an
 94 achievable goal impossible. We also establish operational correspondence through subject
 95 reduction and session fidelity for the direct judgment over a labelled transition system. For
 96 single-label communication, we mechanize the cases where other participants continue to
 97 act concurrently; the world-changing action cases are proved on paper under explicit side
 98 conditions.
- 99 3. **A decidability boundary** (Appendix D). Achievability is decidable when the set of par-
 100 ticipants is fixed. We also show that an extension allowing the dynamic addition of an
 101 unbounded number of participants becomes undecidable.
- 102 4. **An implementation and evaluation** (Sections 4–5). We implement the checker together
 103 with an LLM-compaction front-end and a schema gate, and evaluate it on a corpus covering
 104 the main failure modes. The results are consistent with the theoretical guarantees: we
 105 observe no false impossibility verdicts, while every false achievability verdict falls into one
 106 of the two deferred cases.

107 2 Overview by Example

108 Consider a one-agent skill whose goal is *a flight is booked and a confirmation is sent*. The LLM
 109 compacts the skill into a formal pack. If no email tool is available, there is no action that can establish
 110 confirmation_sent.

$$\begin{aligned}
 \varphi_{goal} &= \text{booked} \wedge \text{confirmation_sent} \\
 \Gamma(a) &= \langle pre_a, \langle Add_a, Del_a, Asg_a, Nd_a \rangle \rangle \\
 \text{search} &\mapsto \langle \top, \langle \{\text{searched}\}, \emptyset, \emptyset, \emptyset \rangle \rangle \\
 \text{filter} &\mapsto \langle \text{searched}, \langle \{\text{filtered}\}, \emptyset, \emptyset, \emptyset \rangle \rangle \\
 \text{book} &\mapsto \langle \text{filtered}, \langle \{\text{booked}\}, \emptyset, \emptyset, \emptyset \rangle \rangle \\
 \text{email} &\mapsto \langle \text{booked}, \langle \{\text{confirmation_sent}\}, \emptyset, \emptyset, \emptyset \rangle \rangle
 \end{aligned}$$

$$\text{dom}(\Gamma_A) = \{\text{search}, \text{filter}, \text{book}\} \quad \text{dom}(\Gamma_B) = \text{dom}(\Gamma_A) \cup \{\text{email}\}$$

Under STRIPS [17] frame semantics, the checker therefore returns *impossible* and identifies the missing capability. If the email tool is available, the goal is reachable through the sequence $\text{search} \cdot \text{filter} \cdot \text{book} \cdot \text{email}$, which the checker returns as a witness. The same example is mechanized in Appendix D (the instance `FlightInstance`). Importantly, the booking itself remains reachable in the first case; the refutation concerns only the missing capability needed to complete the goal.

Tolerance is also important. A skill should not be rejected simply because an agent sends extra status messages, takes a clarifying step, or uses a different payload. Our analysis therefore allows these variations in three ways: *existential may-reachability*, where one successful path is sufficient; *session subtyping*, which allows safe interface variation; and *goal-relevant abstraction*, which ignores payload details that do not affect the goal. These mechanisms are defined in Appendices A–C.

3 Verification Model and Signals

Declared pack. The checker consumes $P = \langle \Gamma, G, \varphi_{\text{goal}}, W_0 \rangle$. Here, Γ contains capability declarations with guards and STRIPS-style effects; G is a goal-marked global protocol (a *global type*; Appendix A); φ_{goal} is the logical goal; and W_0 is the initial abstract world state. The LLM produces P from the natural-language skill. A deterministic schema gate enforces the decidable restrictions: finite static roles, guarded recursion, finite-domain booleans, and quantifier-free linear integer arithmetic. The formal syntax and semantics are given in Appendices A–B.

Verification signals. The checker uses three static signals:

Signal	Role	Output
tool availability and effects	capability reachability	missing capability or witness action
coordination structure	global protocol conformance	protocol mismatch or conformant protocol
goal and cost refinements	logical reachability	blocked guard or model

These signals are checked before execution. Runtime payload checks remain separate obligations and are not assumed by the static checker.

Decision procedure. After the schema gate, the checker (i) rejects undeclared capabilities and goal conditions that cannot be achieved by any available capability; (ii) checks role behaviors against G , requiring exact sender choices while allowing safe receiver-side branch supersets; and (iii) explores existential protocol/world successors from W_0 , using Z3 [18] for guards and refinements. A goal condition that cannot be satisfied blocks the search. If the solver times out or the pack falls outside these restrictions, the checker returns UNKNOWN rather than *impossible*. The checker therefore returns:

IMPOSSIBLE(F)	a sound explanation of why the goal cannot be reached,
ACHIEVABLE(π, O)	a witness path π and deferred obligations O ,
UNKNOWN(r)	an abstention with a reason r .

Each verdict records the pack digest, checker semantics, and artifact version.

Guarantees. Let a configuration consist of a protocol state and an abstract world state, and let *abs* map concrete configurations to abstract configurations. Assume step and goal simulation for the declared pack (Appendix D). If no abstract run reaches φ_{goal} , no concrete run represented by that pack reaches the concrete goal (Theorem 1). The analysis remains sound when the abstraction is made less detailed, and adding capabilities cannot turn an achievable goal into an *impossible* one (Theorems 2–3). Direct conformance satisfies subject reduction and session fidelity (Theorems 4–5). The Coq development covers single-label communication with concurrent actions by other participants; world-changing interleavings require explicit stability and commutativity assumptions. The decidable restrictions above guarantee decidability. Allowing an unbounded dynamic addition of participants is an explicit undecidable extension, for which the implementation returns UNKNOWN. Full rules, assumptions, proofs, and mechanization scope appear in Appendix D.

3.1 Structured feedback and human oversight

The checker returns an explanation for IMPOSSIBLE and a witness path for ACHIEVABLE. These outputs provide concrete feedback rather than a single score. A bounded repair adapter may ask the LLM to fix a protocol, but it cannot invent capabilities or weaken the goal. Every revision is checked again by the deterministic checker. Human review remains necessary for the gap between the original intent and the formal pack. Automated prompt or scaffold search is an application of these verdicts, not an empirical contribution of this paper.

4 Implementation

The executable decision path is approximately 1.45K lines of Python, including the checker, pack gate, formula layer, and session adapter. The compaction front-end issues the artifact prompt to an LLM and passes the result through a deterministic *schema gate* before the checker sees it, so malformed packs are rejected without crashing the trusted core. The Coq artifact contains three developments and two assumption-audit harnesses: reachability soundness core (Theorems 1–3 plus `FlightInstance`, ~230 lines); the direct-typing core (the head-move action safety plus `HandoffInstance`, ~310 lines); and the operational-correspondence core (Theorems 4–5 for the single-label communication fragment with full interleaving, ~250 lines); all three are checked in CI under pinned Coq 8.18.0 with no axioms, verified by `Print Assumptions`. The checker implements the corresponding abstractions from Appendices A–C: STRIPS frame semantics for the world, existential search over the protocol control structure, and Z3 for guard/goal satisfiability and arithmetic refinements. On success it returns the witness path; on refutation, the blocking frontier (missing capability, unsatisfiable guard, unestablished goal conjunct, or non-projectable handoff). The optional LLM front-end may use a `NON_PROJECTABLE` counterexample for one bounded, structure-only repair round; it may not invent capabilities or weaken the goal, and the deterministic checker remains the final referee. This is a small reflective-verification loop, not a general scaffold-search result.

The conformance step ($G \vdash (\mathbb{M}; W)$, Section C.2) is implemented by means of a simple *type inference algorithm* since all typing rules are structural in $(\mathbb{M}; W)$. At each step we freely choose either one or two participant processes to reduce or a goal satisfied by the world obtaining the set of sessions which must be typed as premises of the typing rule.

The adapter also discharges the participant side conditions of T-COMM/T-ACT/T-GOAL: it computes $\text{prt}(G)$ and compares it with the roles the pack declares behaviour for. One direction refutes — a declared role outside $\text{prt}(G)$ has contract `End`, which nothing non-trivial refines. The other cannot: a participant the pack declares nothing for offers no behaviour to refute, yet the judgment of Section C.2 ranges over a whole session with $\text{prt}(G) = \text{prt}(\mathbb{M})$, so the checker returns those roles alongside the verdict rather than leaving the assumption implicit.

5 Evaluation

We assembled a corpus of 15 skills spanning the canonical failure modes — hallucinated planning, retry-forever loops, coordination deadlock, refinement violations — with achievable controls (including detours and informed branches) and two deliberately *spurious* cases exercising the deferred edges. Each skill carries a ground-truth semantic label; the positive class is *achievable*. Evaluation is deterministic and CPU-only: it requires no training or live model inference, completes the 15-case matrix in under one second on a commodity laptop, and gives each individual Z3 query a 10-second timeout.

	predicted achievable	predicted impossible
truly achievable	TP = 6	FN = 0
truly impossible	FP = 2	TN = 7

Two claims, both confirmed. **Soundness:** $\text{FN} = 0$ — no achievable goal was wrongly refuted, the empirical face of Theorem 1 (achievability recall $6/6 = 100\%$). **Incompleteness, contained:** $\text{FP} = 2$, and both are exactly the planted spurious cases — one a false payload post-condition (Layer-C obligation), one an under-specified goal (intent edge). No structural failure was missed in

200 this proof-of-concept corpus; this is evidence of coverage, not a proof that the two residues exhaust
 201 all empirical failure modes. Each refutation reason fires on its matching mode:

	Failure mode (source)	Verdict reason
	hallucinated planning — invokes a non-existent tool	MISSING_CAPABILITY
202	hallucinated planning — no establisher for a goal conjunct	GOAL_UNSAT
	retry-forever loop — guard never satisfiable	BLOCKED_GUARD
	coordination deadlock — unobserved choice / bad handoff	NON_PROJECTABLE
	refinement violation — budget unsatisfiable on every run	GOAL_UNSAT

203 **What the check costs.** With no model in the trusted path, the checker and the deterministic front-
 204 end spend *zero* tokens; only compaction spends any, once per skill *version* and linearly in the skill’s
 205 length, against a run it avoids that recurs per *invocation* and is quadratic in its turns. Under an explicit,
 206 conservative model of that run (Appendix E), compacting the 15-spec corpus costs 12.2K tokens
 207 and avoids $\sim 1.07\text{M}$ per round of invocations — $\sim 87\times$ leverage, unbounded deterministically — so
 208 break-even falls inside the first prevented run in every failure mode, while a skill that is *not* broken
 209 pays 2.9% of one successful run for the check, or nothing.

210 A separate six-case fragment suite exercises behavior absent from the headline matrix: tail-recursive
 211 retry and non-terminating control, dynamic spawning, protocol-independent refutation despite spawn-
 212 ing, and conformant versus non-conformant declared role behavior. It yields two ACHIEVABLE,
 213 three IMPOSSIBLE, and one UNKNOWN. The abstention is reported separately rather than counted
 214 as a false negative, because UNKNOWN makes no refutation claim. Together the suites contain 21
 215 declared packs; the 15-case matrix remains the only binary ground-truth matrix.

216 6 Related Work

217 **Prior-art note.** The claims below reflect the authors’ survey at submission time; the fusion’s novelty
 218 should be confirmed against a systematic search of the planning-meets-behavioural-types literature,
 219 which is dispersed across communities.

220 *Multiparty session types* [13, 16, 9] supply deadlock-freedom via projection and subtyping [14, 10];
 221 our calculus and directed-choice global types follow the modern synchronous formulation of the
 222 Yoshida line [9–11]. We adopt the synthetic line [1, 2]: Section C.2 checks conformance directly
 223 against a whole session configuration rather than projected local types. Our direct rule corresponds
 224 to the conservative, plain-merge condition that bystanders behave uniformly across an unobserved
 225 choice; it does not subsume every protocol accepted by modern full merge. The new contribution is
 226 the world and goal layer: capability guards, effects, goal observation, and the asymmetric refutation
 227 guarantee are not modelled by classical MPST or new SMPS. *Automated planning* [17] supplies
 228 capability-guarded goal reachability; we reuse its frame semantics but embed it in a multiparty
 229 choreography. The contribution is the *fusion*, decided over an goal-marked global type synthesized by
 230 an LLM. Recent agent-verification work is adjacent but differently aimed: *provable coordination via*
 231 *message sequence charts* [20] projects a global choreography for LLM agents but targets coordination
 232 correctness, not goal achievability; its runtime-planning extension also establishes that LLM synthesis
 233 of a coordination protocol is itself prior art. *VeriGuard* [21] separates offline policy verification
 234 from online monitoring as a safety shield; *agent behavioural contracts* [22] provide probabilistic,
 235 runtime-enforced compliance and composition conditions for multi-agent chains, whereas our verdict
 236 is static and sound for refutation relative to the pack.

237 *Skill-bundle scanners and verified-skill catalogs* are complementary. NVIDIA’s SkillSpector [23] and
 238 Verified Agent Skills pipeline [24] address the same deployment object—portable agent skills—but
 239 ask a different question. SkillSpector is a pre-publication security control: it normalizes a skill bundle
 240 and runs deterministic scanners such as static pattern detectors, AST and taint analyzers, manifest-
 241 consistency checks, metadata-poisoning checks, supply-chain checks, and optional LLM-backed
 242 semantic analyzers. This is valuable admission control for skills, but it is not a compiler or type
 243 discipline in the sense used here. As published, SkillSpector does not aim to provide an operational
 244 semantics of goal-marked skills, a goal-reachability decision, or a refutation-soundness theorem. A
 245 SkillSpector-like pass belongs at the boundary of the compiler: before or alongside LLM compaction,
 246 it can vet the incoming SKILL.md, code, dependencies, and MCP metadata; after compaction, it
 247 can help justify the honesty of declared capabilities and least-privilege metadata. The type checker

248 then consumes the sanitized formal pack and proves the narrower claim that the declared goal is not
249 structurally impossible.

250 None of the above frames the static check as goal achievability over a declared world and goal-marked
251 protocol, nor gives the same refutation-sound/achievement-incomplete asymmetry.

252 7 Limitations and Future Work

- 253 • **Relative soundness.** Everything is proved about the declared Γ and S . If the prose lies (a
254 “read-only” tool that writes), the checker verifies a fiction; honest-declaration is a separate
255 obligation, of which this makes only the interaction slice decidable. A practical compiler
256 should therefore include a skill-bundle security pass—for example a SkillSpector-like
257 scanner—to inspect SKILL.md, scripts, dependencies, manifests, and permission metadata
258 before the formal pack is trusted by the achievability checker.
- 259 • **Static topology for totality.** Dynamic spawning drops the procedure to a semi-decision
260 (Proposition 2); a practical system answers UNKNOWN rather than guessing.
- 261 • **Trusted abstraction.** A coarse α silently widens tolerance (more false achievable), never
262 breaking Theorem 1 but weakening usefulness.
- 263 • **Engineering gaps.** The mechanization covers the generic refutation-sound abstraction
264 theorem and concrete examples, head-move action safety, and subject reduction/session
265 fidelity with full bystander interleaving for the single-label communication model, all axiom-
266 free. The executable checker is schema-gated and tested against that specification, but the
267 STEP-SIM/GOAL-SIM instantiation for its exact symbolic state is not mechanized. The Coq
268 correspondence model also does not yet combine observable goal labels, participant match-
269 ing, multi-branch choice, and world-changing interleaving in one development. Still open
270 are that faithful mechanization; a head-normalization/permutation bridge from interleaved
271 session runs to the head-ordered reachability search; equivalence between the declarative
272 judgment and the projection-based adapter; and a concrete-session realization theorem for
273 abstract witnesses. The 21-pack evaluation remains a proof of concept.

274 **Broader impact.** Pre-execution refutation reduces wasted computation (Appendix E) and prevents
275 structurally impossible agent plans from reaching deployment. The same formal verdict can, however,
276 create false confidence when an inaccurate pack omits a harmful tool effect or weakens the user’s
277 actual goal. We therefore expose the declared pack, witness or blocking frontier, and deferred
278 obligations rather than presenting a single trust score. Intent review and runtime payload monitoring
279 remain necessary before the method is used in consequential settings; this work does not evaluate
280 fairness, privacy, or deployment safety.

281 8 Conclusion

282 An LLM can draft an agent’s plan from intent; the open question has been whether anything can tell,
283 cheaply and before execution, that the plan is doomed. For the *achievability* question the answer is
284 yes, relative to a declared pack: combine planning’s capability-guarded reachability with a synthetic,
285 directly-typed global protocol; treat LLM compaction as untrusted; and use a refutation-sound
286 abstraction so an *impossible* verdict is never wrong about that declared model. *Achievable* remains
287 honest about payload and intent fidelity, while *Unknown* abstains outside static topology. This yields
288 a structured, inspectable verification signal for agent skills and a path toward, rather than a completed
289 proof of, stronger per-step safety.

290 References

- 291 [1] Franco Barbanera, Mariangiola Dezani-Ciancaglini & Ugo de’Liguoro (2024): *Un-projectable Global*
292 *Types for Multiparty Sessions*. In Alessandro Bruni, Alberto Momigliano, Matteo Pradella & Matteo Rossi,
293 editors: *PPDP*, ACM Press, pp. 15:1–15:13,
- 294 [2] Francesco Dagnino, Paola Giannini & Mariangiola Dezani-Ciancaglini (2023): *Deconfined Global Types*
295 *for Asynchronous Sessions*. *Logical Methods in Computer Science* 19(1), pp. 1–41,

- [3] J. Li, X. Xiao, Y. Zhang, C. Liu, L. Zhao, X. Liao, Y. Ji, J. Wang, J. Gu, Y. Ge, W. Xu, X. Fang, X. Xu, T. Zhao, Y. Kim, T. Wang, J. Hamm, S. Krishnaswamy, J. Huan, C. Reddy. Agent Harness Engineering: A Survey. 2026.
- [4] M. Chen, J. Wang, Z. Liu, Y. Wang, H. Zheng, Q. Wang. From Failed Trajectories to Reliable LLM Agents: Diagnosing and Repairing Harness Flaws. arXiv:2606.06324, 2026.
- [5] Agent Skills Standard. Specification. <https://github.com/agentskills/agentskills/blob/main/docs/specification.mdx>, 2025.
- [6] S. Hu, C. Lu, J. Clune. Automated Design of Agentic Systems. *ICLR*, 2025.
- [7] J. Zhang, J. Xiang, Z. Yu, F. Teng, X. Chen, J. Chen, M. Zhuge, X. Cheng, S. Hong, J. Wang, et al. AFlow: Automating Agentic Workflow Generation. *ICLR*, 2025.
- [8] M. Cemri, M. Z. Pan, S. Yang, L. A. Agrawal, B. Chopra, R. Tiwari, K. Keutzer, A. Parameswaran, D. Klein, K. Ramchandran, M. Zaharia, J. E. Gonzalez, I. Stoica. Why Do Multi-Agent LLM Systems Fail? arXiv:2503.13657, 2025.
- [9] N. Yoshida, L. Gheri. A Very Gentle Introduction to Multiparty Session Types. *ICDCIT*, LNCS 11969, 2020.
- [10] S. Ghilezan, S. Jakšić, J. Pantović, A. Scalas, N. Yoshida. Precise Subtyping for Synchronous Multiparty Sessions. *J. Log. Algebr. Meth. Program.* 104, 2019.
- [11] A. Scalas, N. Yoshida. Less is More: Multiparty Session Types Revisited. *Proc. ACM Program. Lang.* 3(POPL), 2019.
- [12] K. Honda, V. T. Vasconcelos, M. Kubo. Language Primitives and Type Discipline for Structured Communication-Based Programming. *ESOP*, LNCS 1381, pp. 122–138, 1998.
- [13] K. Honda, N. Yoshida, M. Carbone. Multiparty Asynchronous Session Types. *J. ACM* 63(1), 2016.
- [14] S. Gay, M. Hole. Subtyping for Session Types in the Pi Calculus. *Acta Informatica* 42(2–3), 2005.
- [15] D. Brand, P. Zafiropulo. On Communicating Finite-State Machines. *J. ACM* 30(2), 1983.
- [16] E. Li, F. Stutz, T. Wies, D. Zufferey. Complete Multiparty Session Type Projection with Automata. *CAV* 2023.
- [17] R. Fikes, N. Nilsson. STRIPS: A New Approach to the Application of Theorem Proving to Problem Solving. *Artif. Intell.* 2(3–4), 1971.
- [18] L. de Moura, N. Bjørner. Z3: An Efficient SMT Solver. *TACAS*, LNCS 4963, pp. 337–340, Springer, 2008.
- [19] C. Barrett, P. Fontaine, C. Tinelli. The SMT-LIB Standard: Version 2.6. Technical Report, Department of Computer Science, The University of Iowa, 2017. <https://smt-lib.org>.
- [20] B. Bollig, M. Függer, T. Nowak. Provable Coordination for LLM Agents via Message Sequence Charts. *ISoLA*, 2026. arXiv:2604.17612.
- [21] L. Miculicich, M. Parmar, H. Palangi, K. D. Dvijotham, M. Montanari, T. Pfister, L. T. Le. VeriGuard: Enhancing LLM Agent Safety via Verified Code Generation. arXiv:2510.05156, 2025.
- [22] V. P. Bhardwaj. Agent Behavioral Contracts: Formal Specification and Runtime Enforcement for Reliable Autonomous AI Agents. arXiv:2602.22302, 2026.
- [23] N. Paz, K. Pradeep, N. Raghavan, A. Nikirk, M. Gupta. SkillSpector: A Pre-Publication Security Control for Agent Skills. *Agent Skills Workshop at CAIS*, 2026.
- [24] NVIDIA. NVIDIA-Verified Agent Skills Provide Capability Governance for AI Agents. NVIDIA Technical Blog, 19 May 2026.

337 A Syntax

338 We use disjoint sets of predicates $p \in \text{Pred}$ (boolean), numeric variables $x \in \text{Var}$, roles $p, q, r \in \mathcal{R}$,
 339 capability names $a \in \text{Cap}$, and message labels $\ell \in \text{Lab}$. Following the modern presentation of
 340 multiparty **sessions** [1, 2] we define the model from the bottom up: state and capabilities (Sections A.1–
 341 A.2), then *processes* (Section A.3), and finally the global *type* they are checked against directly
 342 (Section A.4).

343 In each grammar, “ $::=$ ” lists the possible forms of an expression. In each inference rule below,
 344 the statement below the line holds when all premises above the line hold. Rule names identify the
 345 corresponding checker obligation; they are not additional assumptions.

346 A.1 State logic

$$\begin{aligned} e &::= x \mid n \mid e + e \mid e - e \mid c \cdot e & (c, n \in \mathbb{Z}) \\ \varphi &::= p \mid \top \mid \perp \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid e \bowtie e & \bowtie \in \{<, \leq, =, \geq, >, \neq\} \end{aligned}$$

347 The state logic $\mathcal{L}$ is quantifier-free linear integer arithmetic [19] with uninterpreted boolean predicates
 348 — a decidable fragment for which entailment $\models$ and satisfiability are discharged by a satisfiability-
 349 modulo-theories (SMT) solver.

350 A.2 Worlds and capabilities

351 A *world* $W = \langle B, N \rangle$ assigns truth values to predicates $B : \text{Pred} \rightarrow \mathbb{B}$ and numeric variables
 352 $N : \text{Var} \rightarrow \mathbb{Z}$. A *capability context* Γ maps each available tool a to a guarded effect:

$$\begin{aligned} \Gamma(a) &= \langle \text{pre}_a, \text{eff}_a \rangle, & \text{pre}_a &\in \mathcal{L}, \\ \text{eff}_a &= \langle \text{Add}_a \subseteq \text{Pred}, \text{Del}_a \subseteq \text{Pred}, \text{Asg}_a : \text{Var} \rightarrow e, \text{Nd}_a : \text{Var} \rightarrow \mathcal{L} \rangle \end{aligned}$$

353 $a \notin \text{dom}(\Gamma) \iff$ the agent does *not* possess tool a (no hallucinated tools).

354 *Add/Del* are the STRIPS add- and delete-lists; *Asg* gives deterministic numeric updates $x := e$; *Nd*
 355 gives *nondeterministic* updates $x := *$ constrained by a formula over the new value (used to model
 356 post-conditions such as “*books a fare under 500*”).

357 A.3 Processes and multiparty sessions

358 Skills *execute* as processes; a multi-agent run is a session of named processes. We use the synchronous
 359 multiparty session calculus in its modern coinductive **presentation** [1, 2]: **processes coinductively**
 360 **defined are possibly infinite regular trees (finitely many distinct subtrees)**, and labels are goal-relevant
 361 tokens, $\ell = \alpha(m)$ for concrete messages m — the tolerance dial α is a parameter of the label
 362 alphabet.

$$P ::=^{\text{coind}} \mathbf{0} \mid \mathbf{p}!\{\ell_i.P_i\}_{i \in I} \mid \mathbf{p}?\{\ell_i.P_i\}_{i \in I} \mid a.P \quad \mathbb{M} = \mathbf{p}_1[P_1] \parallel \cdots \parallel \mathbf{p}_n[P_n]$$

363 with I finite and non-empty, labels ℓ_i pairwise distinct, and **role** names pairwise distinct. The output
 364 $\mathbf{p}!\{\ell_i.P_i\}$ *internally* chooses a label and sends it to $\mathbf{p}$; the input $\mathbf{p}?\{\ell_i.P_i\}$ offers an *external* choice;
 365 $a.P$ invokes capability a — the only construct that changes the world. A skill S_p is a process: the
 366 behaviour declared for role $\mathbf{p}$.

367 A.4 Global types

368 The single communication construct of a global type is a *directed* labelled choice — selection by $\mathbf{p}$
 369 and branching at $\mathbf{q}$ in one construct. (A partnerless “ $\mathbf{p}$ selects” admits no coherent view for the roles
 370 that did not choose; directed choice is the standard MPST construct [13, 9].)

$$G ::=^{\text{coind}} \mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G_i\}_{i \in I} \mid a@\mathbf{p}.G \mid \check{\varphi}.G \mid \text{End}$$

371 subject to $\mathbf{p} \neq \mathbf{q}$ and the same regularity and label conditions as processes. This is the *only* type
 372 in the discipline: Section C.2 types a whole session directly against G — there is no local-type
 373 grammar, no projection function, and no separate subtyping relation. Goal markers $\check{\varphi}$ live *only*
 374 in G : achievability is decided on the global type (Rule ACH in C.3), and markers are consumed at
 375 typing time by T-GOAL without ever burdening a process. Only $a@\mathbf{p}$ changes the world; messages
 376 carry coordination information only.

B Operational Semantics

B.1 World transitions

A capability fires when its guard holds, producing a successor world. Because Nd is nondeterministic, $\langle W \rangle a \langle W' \rangle$ may relate one world to many. The relation is implicitly indexed by Γ : the notation pre_a, eff_a is defined only when $\Gamma(a) = \langle pre_a, eff_a \rangle$, so every firing presupposes $a \in \text{dom}(\Gamma)$.

$$\frac{\text{WORLD-ACT} \quad W \models pre_a \quad W' \in \text{apply}(eff_a, W)}{\langle W \rangle a \langle W' \rangle}$$

$$\text{apply}(eff_a, \langle B, N \rangle) = \langle B', N' \rangle, \quad B' = (B \cup Add_a) \setminus Del_a$$

$$N'(x) = \begin{cases} \llbracket Asg_a(x) \rrbracket_N & \text{if } x \in \text{dom}(Asg_a) \\ v \text{ with } \langle B, N[x \mapsto v] \rangle \models Nd_a(x) & \text{if } x \in \text{dom}(Nd_a) \\ N(x) & \text{otherwise (frame)} \end{cases}$$

B.2 A labelled semantics for sessions and global types

Only capabilities touch the world, so sessions reduce as configurations $(\mathbb{M}; W)$; communication is synchronous [10]. Each transition carries a label Λ recording its *participants*, so that independent moves by disjoint roles can be identified. Two auxiliary maps drive the labelled system: the *participants* $\text{prt}(\cdot)$ of a global type, a label, or a session, and the *capabilities* $\text{cap}(\cdot)$ (the set of moves a global type can perform). On global types,

$$\begin{aligned} \text{prt}(p \rightarrow q : \{\ell_i. G_i\}_{i \in I}) &= \{p, q\} \cup \bigcup_{i \in I} \text{prt}(G_i) & \text{prt}(\check{\varphi}.G) &= \text{prt}(G) \\ \text{prt}(a @ p.G) &= \{p\} \cup \text{prt}(G) & \text{prt}(\text{End}) &= \emptyset, \end{aligned}$$

on labels $\text{prt}(p \ell q) = \{p, q\}$, $\text{prt}(a @ p) = \{p\}$, $\text{prt}(\check{\varphi}) = \emptyset$, and on sessions $\text{prt}(\mathbb{M}) = \{p \mid \mathbb{M} \equiv p[P] \ \& \ P \neq 0\}$ (the roles with a residual behaviour). The capabilities of a global type are

$$\begin{aligned} \text{cap}(p \rightarrow q : \{\ell_i. G_i\}_{i \in I}) &= \{p \ell_i q \mid i \in I\} \cup \bigcup_{i \in I} \text{cap}(G_i) \\ \text{cap}(a @ p.G) &= \{a @ p\} \cup \text{cap}(G) & \text{cap}(\check{\varphi}.G) &= \{\check{\varphi}\} \cup \text{cap}(G) & \text{cap}(\text{End}) &= \emptyset. \end{aligned}$$

The capability names used by a protocol are

$$\text{caps}(G) = \{a \mid \exists p. a @ p \in \text{cap}(G)\}.$$

The LTS of sessions is defined by

$$\begin{aligned} \text{S-COMM} \quad & \frac{k \in I \subseteq J}{(p[q!\{\ell_i.P_i\}_{i \in I}] \parallel q[p?\{\ell_j.Q_j\}_{j \in J}] \parallel \mathbb{M}; W) \xrightarrow{p \ell_k q} (p[P_k] \parallel q[Q_k] \parallel \mathbb{M}; W)} \\ \text{S-ACT} \quad & \frac{\langle W \rangle a \langle W' \rangle}{(p[a.P] \parallel \mathbb{M}; W) \xrightarrow{a @ p} (p[P] \parallel \mathbb{M}; W')} \\ \text{S-GOAL} \quad & \frac{W \models \varphi}{(\mathbb{M}; W) \xrightarrow{\check{\varphi}} (\mathbb{M}; W)} \end{aligned}$$

where $\parallel$ is commutative, associative and $p[0] \parallel \mathbb{M} \equiv \mathbb{M}$. A non-terminated configuration with no successor is *stuck*; the deadlocking planner/worker handoff of Section 1 is the stuck two-role configuration in which both processes begin with an input. A goal marker is *not* carried by any process; instead S-GOAL lets a configuration *observe* that its world already entails one of its declared goals φ , emitting the label $\check{\varphi}$ without altering $\mathbb{M}$ or W . (Each session carries a set of goal formulas, and φ ranges over that set — not over every formula the world happens to satisfy — so that each observation has a matching marker in the protocol.) This observable goal step is matched at the type level by G-GOAL-E below, so that $\check{\varphi}$ is a genuine label Λ of both transition systems and the operational correspondence of Section D.3 ranges over it uniformly.

The checker relates a session to its protocol through a *labelled* transition system on global types, $(G, W) \rightsquigarrow_{\Lambda} (G', W')$, built from *head* rules that consume the leading construct and *interleaving*

rules that let a move by participants *disjoint* from the head commute inward:

$$\begin{array}{c}
\text{G-COMM-E} \quad \frac{k \in I}{(\mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G_i\}_{i \in I}, W) \rightsquigarrow_{\mathbf{p} \ell_k \mathbf{q}} (G_k, W)} \quad \text{G-ACT-E} \quad \frac{\langle W \rangle a \langle W' \rangle}{(a @ \mathbf{p}. G, W) \rightsquigarrow_{a @ \mathbf{p}} (G, W')} \\
\\
\text{G-GOAL-E} \quad \frac{W \models \varphi}{(\check{\varphi}. G, W) \rightsquigarrow_{\check{\varphi}} (G, W)} \\
\\
\text{G-COMM-I} \quad \frac{(G_i, W) \rightsquigarrow_{\Lambda} (G'_i, W') \quad \Lambda \in \text{cap}(G_i) \quad \forall i \in I \quad \text{prt}(\Lambda) \cap \{\mathbf{p}, \mathbf{q}\} = \emptyset}{(\mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G_i\}_{i \in I}, W) \rightsquigarrow_{\Lambda} (\mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G'_i\}_{i \in I}, W')} \\
\\
\text{G-ACT-I} \quad \frac{(G, W) \rightsquigarrow_{\Lambda} (G', W') \quad \Lambda \in \text{cap}(G) \quad \text{prt}(\Lambda) \cap \{\mathbf{p}\} = \emptyset}{(a @ \mathbf{p}. G, W) \rightsquigarrow_{\Lambda} (a @ \mathbf{p}. G', W')} \\
\\
\text{G-GOAL-I} \quad \frac{(G, W) \rightsquigarrow_{\Lambda} (G', W') \quad \Lambda \in \text{cap}(G)}{(\check{\varphi}. G, W) \rightsquigarrow_{\Lambda} (\check{\varphi}. G', W')}
\end{array}$$

399 G-GOAL-E discharges a goal marker when $W \models \varphi$, emitting the observable label $\check{\varphi}$ that matches
400 S-GOAL. In the head-only reduction used for achievability, an unsatisfied marker is a checkpoint and
401 blocks the path. In the full labelled correspondence, G-GOAL-I lets a continuation step commute
402 inward under a still-pending marker. When that step changes the world, the correspondence theorem
403 uses the goal-stability hypothesis stated in Section D.3; removing that hypothesis is future work. The
404 extraction (head) rules G-COMM-E/G-ACT-E/G-GOAL-E give the Λ -erased reduction $(G, W) \rightsquigarrow$
405 (G', W') that the achievability judgment below uses; the interleaving rules matter only for the
406 operational correspondence of Section D.3, and their $\Lambda \in \text{cap}(G)$ premise confines a commuting
407 move to a capability the residual protocol actually offers. G-COMM-E is legitimate because choice
408 is directed — the branch is a communication both endpoints observe, and Section C.2 checks that
409 every third party can follow. A well-formed pack uses its declared goal φ_{goal} at each goal marker. A
410 configuration is *goal-satisfying* when control is at such a marker or at the terminus and the world
411 entails that declared goal:

$$\begin{array}{c}
\text{goal}_{\varphi_{\text{goal}}}(\check{\varphi}_{\varphi_{\text{goal}}}.G, W) \iff W \models \varphi_{\text{goal}}, \quad \text{goal}_{\varphi_{\text{goal}}}(\text{End}, W) \iff W \models \varphi_{\text{goal}}, \\
412 \quad \Gamma; G \models \Diamond_{\text{init}} \varphi_{\text{goal}} \iff \exists W_0 \models \text{init}. \exists (G', W'). (G, W_0) \rightsquigarrow^* (G', W') \wedge \text{goal}_{\varphi_{\text{goal}}}(G', W').
\end{array}$$

413 The diamond is *existential*: a single witnessing run suffices — the formal source of detour-tolerance.

414 B.3 Realizability and subtyping, without projection

415 The classical route to ruling out a deadlocking handoff is *projection*: compute each role's local
416 type $G|_r$ by structural recursion on G , taking the *merge* of a choice's branches at every role that
417 neither sends nor receives, and declaring G *realizable* exactly when this partial function is defined
418 everywhere — undefinedness of the merge is the static signature of a role forced to behave differently
419 in branches it cannot distinguish. Session subtyping $\leq$ is then a second, separate coinductive relation
420 on local types, letting a role's actual behaviour offer more external choices and fewer internal ones
421 than its projected contract demands [14, 10, 11].

422 Both devices exist to answer questions about the *network as a whole* — “can every role tell the
423 branches apart it needs to?”, “may this implementation stand in for that contract?” — while type-
424 checking only one role's syntax at a time. Section C.2 answers the same two questions with a single
425 judgment, *typing the whole configuration directly against G*: a choice node's rule quantifies over
426 every branch the protocol declares and requires the *same* residual configuration to check against each
427 one. This is the conservative plain-merge condition, obtained without ever computing $\sqcap$ or a local
428 type. The receiver-side subtyping direction (a receiver may accept a superset of the protocol's labels)
429 is folded into the same rule as a label-set inclusion $I \subseteq J$. Tolerance therefore still has the same
430 three independent knobs from Section 2 — the existential $\Diamond$ (Section B.2), subtyping (now inside

431 T-COMM, Section C.2), and the granularity of α — “flexible, tolerant of small details” is still the
 432 profile $\langle \text{coarse } \alpha, \Diamond, \text{linear } \mathcal{L} \rangle$; a later safety layer flips the same machinery to the universal $\Box$ over
 433 permission predicates.

434 C The Achievability Type System

435 Achievability now factors into three premises, each separately decidable, combined by the top rule
 436 ACH: no hallucinated capability, direct conformance of the declared skills to G , and reachability of
 437 the goal.

438 C.1 Capability and goal typing

$$\begin{array}{c} \text{T-EFF} \\ \frac{\Gamma(a) = \langle pre, eff \rangle \quad \varphi \models pre}{\Gamma \vdash \{\varphi\} a \{\text{sp}(\varphi, eff)\}} \end{array} \quad \begin{array}{c} \text{T-MISS} \\ \frac{a \notin \text{dom}(\Gamma)}{\Gamma \vdash a \triangleright \mathbf{X} \text{ (MISSING_CAPABILITY)}} \end{array}$$

439 $\text{sp}(\varphi, eff)$ is the strongest postcondition of the STRIPS effect. T-MISS fires on hallucinated planning:
 440 the plan names a tool the context does not grant, and achievability is refuted without any search.

441 C.2 Direct conformance typing

Skills are typed *directly* against the global protocol and the world, by the coinductive judgment $G \vdash (\mathbb{M}; W)$ — read “ G types session $\mathbb{M}$ starting from world W .” There is no local type, no projection, and no separate subtyping relation: the receiver-side subtyping direction and the unobserved-choice check are structural side conditions of one rule.

$$\begin{array}{c} \text{T-END} \\ \frac{}{\text{End} \vdash (\mathbf{p}[0]; W)} \end{array} \quad \begin{array}{c} \text{T-ACT} \\ \frac{G \vdash (\mathbf{p}[P] \parallel \mathbb{M}; W') \quad \langle W \rangle a \langle W' \rangle \quad \text{prt}(G) \cup \{\mathbf{p}\} = \text{prt}(\mathbb{M}) \cup \{\mathbf{p}\}}{a @ \mathbf{p}. G \vdash (\mathbf{p}[a.P] \parallel \mathbb{M}; W)} \end{array}$$

$$\begin{array}{c} \text{T-GOAL} \\ \frac{G \vdash (\mathbb{M}; W) \quad W \models \varphi \quad \text{prt}(G) = \text{prt}(\mathbb{M})}{\checkmark \varphi. G \vdash (\mathbb{M}; W)} \end{array}$$

$$\begin{array}{c} \text{T-COMM} \\ \frac{G_i \vdash (\mathbf{p}[P_i] \parallel \mathbf{q}[Q_i] \parallel \mathbb{M}; W) \quad \forall i \in I \subseteq J \quad \text{prt}(\mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G_i\}_{i \in I}) = \text{prt}(\mathbb{M}) \cup \{\mathbf{p}, \mathbf{q}\}}{\mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G_i\}_{i \in I} \vdash (\mathbf{p}[\mathbf{q}!\{\ell_i.P_i\}_{i \in I}] \parallel \mathbf{q}[\mathbf{p}?\{\ell_j.Q_j\}_{j \in J}] \parallel \mathbb{M}; W)} \end{array}$$

442 T-ACT pairs type and world in lockstep with WORLD-ACT: the unconsumed type $a @ \mathbf{p}.G$ sits with
 443 the *pre*-effect world W and the residual G with the *post*-effect world W' — the same convention
 444 as G-ACT-E (Section B.2). In T-COMM the sender $\mathbf{p}$ offers exactly the protocol’s branch set I ,
 445 while the receiver $\mathbf{q}$ may accept *more* ($I \subseteq J$) — the one subtyping direction the rule keeps (safe
 446 interface slack at the receiver). Requiring *every* protocol branch ℓ_i to be checked against the *same*
 447 bystanders $\mathbb{M}$ is the plain-merge form of the unobserved-choice condition, so a role forced to behave
 448 differently across branches it cannot distinguish makes no derivation possible — the rule fails closed,
 449 and deadlock-freedom falls out of T-COMM itself with no merge to compute (Theorem 4). Each
 450 rule additionally carries a *participant-matching* side condition tying the protocol’s participants to
 451 the session’s: $\text{prt}(\mathbf{p} \rightarrow \mathbf{q} : \{\ell_i.G_i\}) = \text{prt}(\mathbb{M}) \cup \{\mathbf{p}, \mathbf{q}\}$ in T-COMM, $\text{prt}(G) \cup \{\mathbf{p}\} = \text{prt}(\mathbb{M}) \cup \{\mathbf{p}\}$
 452 in T-ACT, and $\text{prt}(G) = \text{prt}(\mathbb{M})$ in T-GOAL. These make the invariant “the roles the protocol still
 453 mentions are exactly the roles still running” hold at every node, ruling out a session that carries a stray
 454 participant the protocol has forgotten (or vice versa) — the well-formedness that lets the interleaving
 455 cases of Section D.3 match a bystander move on the session to a $\text{cap}(\cdot)$ -labelled move on the type.
 456 T-END closes the discipline at the terminus: the terminated protocol End types a single idle role $\mathbf{p}[0]$
 457 — and, through the congruence $\mathbf{p}[0] \parallel \mathbb{M} \equiv \mathbb{M}$, any session all of whose roles have terminated —
 458 the base case of the participant invariant, since $\text{prt}(\text{End}) = \emptyset$ and an idle role contributes nothing to
 459 $\text{prt}(\mathbb{M})$. T-ACT is derivable only when $\langle W \rangle a \langle W' \rangle$ holds, which itself presupposes $a \in \text{dom}(\Gamma)$:
 460 this is where hallucinated capabilities die, with T-MISS as the diagnostic.

461 C.3 The achievability judgment

$$\text{ACH} \quad \frac{\text{caps}(G) \subseteq \text{dom}(\Gamma) \quad \exists W_0 \models \text{init}. G \vdash (p_1[S_{p_1}] \parallel \dots \parallel p_n[S_{p_n}]; W_0) \wedge \Gamma; (G, W_0) \models \Diamond \varphi_{\text{goal}}}{\Gamma \vdash \{S_p\}_p : G \triangleright \Diamond \varphi_{\text{goal}}}$$

462 Read top-to-bottom: no hallucinated tools; some plausible initial world exists from which the declared
 463 skills directly conform to the protocol (deadlock-free, every capability call guarded) *and* that same
 464 world reaches the goal. The reachability half is unchanged from Section B.2 and decided on Γ, G
 465 alone, before any S_p is known — exactly what lets the checker run on the LLM-compacted pack
 466 the moment it exists; conformance is the additional, separately decidable check once the skills are
 467 declared, and needs no transport lemma: the receiver-side subtyping slack is already inside T-COMM.

468 C.4 Reachability rules

The diamond is derived by the following inductive system, realized by the checker as a finite forward search over $\rightsquigarrow$; unlike T-COMM/T-ACT/ T-GOAL, it never mentions a session $\mathbb{M}$, which is what makes it decidable on the pack alone.

$$\begin{array}{c} \text{REACH-GOAL} \\ \frac{W \models \varphi_{\text{goal}} \quad G = \text{End} \vee \exists G'. G = \checkmark \varphi_{\text{goal}}. G'}{\Gamma; (G, W) \models \Diamond \varphi_{\text{goal}}} \\[10pt] \text{REACH-STEP} \quad \frac{(G, W) \rightsquigarrow (G', W') \quad \Gamma; (G', W') \models \Diamond \varphi_{\text{goal}}}{\Gamma; (G, W) \models \Diamond \varphi_{\text{goal}}} \quad \text{REACH-OR} \quad \frac{\exists k \in I. \Gamma; (G_k, W) \models \Diamond \varphi_{\text{goal}}}{\Gamma; (p \rightarrow q: \{\ell_i.G_i\}_{i \in I}, W) \models \Diamond \varphi_{\text{goal}}} \end{array}$$

469 D Metatheory

470 We separate what is *mechanized in Coq* (Theorems 1–3; Theorems 4–5 for the single-label com-
 471 munication fragment with full interleaving; the head-move action case; and the concrete instances,
 472 all axiom-free) from what is *proved on paper* (the action-interleaving cases, decidability, and the
 473 undecidability boundary).

474 D.1 Soundness of the abstraction

475 Let $\mathcal{C}$ be the concrete configuration space of protocol/world pairs, $\mathcal{A}$ the abstract configuration space
 476 the checker explores, and $\text{abs} : \mathcal{C} \rightarrow \mathcal{A}$ the abstraction. We require:

$$(\text{step-sim}) \quad C \rightsquigarrow C' \Rightarrow \text{abs}(C) \rightsquigarrow^* \text{abs}(C'), \quad (\text{goal-sim}) \quad \text{goal}_{\varphi_{\text{goal}}}(C) \Rightarrow \hat{\text{goal}}(\text{abs}(C)).$$

477 **Lemma 1** (Transport). *If $C_0 \rightsquigarrow^* C$ then $\text{abs}(C_0) \rightsquigarrow^* \text{abs}(C)$.*

478 *Proof.* By induction on the reflexive-transitive closure. The empty run is reflexivity. For $C_0 \rightsquigarrow^* C'$
 479 $C' \rightsquigarrow C$, the induction hypothesis gives $\text{abs}(C_0) \rightsquigarrow^* \text{abs}(C')$ and STEP-SIM gives $\text{abs}(C') \rightsquigarrow^* \text{abs}(C)$;
 480 compose them. Mechanized as `reach_abs`, parameterized by STEP-SIM. $\square$

481 **Lemma 2** (Reachability monotonicity). *If every step of $\rightsquigarrow_1$ is a step of $\rightsquigarrow_2$ (that is, $x \rightsquigarrow_1 y$ implies
 482 $x \rightsquigarrow_2 y$), then $s \rightsquigarrow_1^* w$ implies $s \rightsquigarrow_2^* w$.*

483 *Proof.* By induction on the closure $s \rightsquigarrow_1^* w$. The empty run transports by reflexivity. For $s \rightsquigarrow_1^* u \rightsquigarrow_1 w$,
 484 the induction hypothesis gives $s \rightsquigarrow_2^* u$; the hypothesis turns the last step $u \rightsquigarrow_1 w$ into $u \rightsquigarrow_2 w$;
 485 appending it gives $s \rightsquigarrow_2^* w$. Mechanized as `reach_mono`. $\square$

486 D.2 Refutation soundness, tolerance, monotonicity

487 **Theorem 1** (Refutation soundness). *If the abstract system has no reachable goal state from
 488 $\text{abs}(G, W_0)$, then no declared run of that pack configuration reaches φ_{goal} . Hence an impossi-
 489 ble verdict is never wrong relative to the pack, provided STEP-SIM and GOAL-SIM hold.*

490 *Proof.* We prove the contrapositive: if *some* declared run reaches the goal, then the abstract
 491 system reaches an abstract goal state. Suppose $C_0 \rightsquigarrow^* C$ with $\text{goal}_{\varphi_{\text{goal}}}(C)$. By Lemma 1,
 492 $\text{abs}(C_0) \rightsquigarrow^* \text{abs}(C)$, and by GOAL-SIM, $\hat{\text{goal}}(\text{abs}(C))$. Thus $\text{abs}(C)$ is an abstract goal state reach-
 493 able from $\text{abs}(C_0)$, contradicting the hypothesis. Hence no declared run reaches φ_{goal} : a verdict
 494 of *impossible* — read off exactly the premise, that the abstract search exhausts without hitting an
 495 abstract goal — is never wrong. Mechanized, axiom-free as a theorem schema parameterized by the
 496 two simulation hypotheses (the existential witness is $\text{abs}(C)$): $\square$

```

497   Theorem refutation_sound : forall s0,
498     (~ exists a, reach astep (abs s0) a /\ agoal a) ->
499     (~ exists w, reach cstep s0 w /\ cgoal w).
500   (* Print Assumptions refutation_sound. => Closed under the global
501     context *)

```

502 **Theorem 2** (Tolerance soundness). *If a coarser over-approximation (more edges) cannot reach the*
 503 *goal, neither can any finer system. Hence coarsening α to tolerate more detail never produces a false*
 504 *refutation.*

505 *Proof.* Let $\rightsquigarrow_{\text{fine}}$ and $\rightsquigarrow_{\text{coarse}}$ be the two abstract step relations, with the coarsening an *over-*
 506 *approximation*: every fine step is also a coarse step, $x \rightsquigarrow_{\text{fine}} y \Rightarrow x \rightsquigarrow_{\text{coarse}} y$. Suppose, for
 507 contradiction, that the finer system reaches the goal: $s_0 \rightsquigarrow_{\text{fine}}^* a$ with $\hat{\text{goal}}(a)$. By Lemma 2 (with
 508 $\rightsquigarrow_1 = \rightsquigarrow_{\text{fine}}$, $\rightsquigarrow_2 = \rightsquigarrow_{\text{coarse}}$), $s_0 \rightsquigarrow_{\text{coarse}}^* a$; and $\hat{\text{goal}}(a)$ still holds. So the coarser system reaches
 509 a goal state, contradicting the hypothesis. Hence no coarser-refuted goal is reachable in any finer
 510 system: coarsening α only *adds* abstract edges, so it can never turn an unreachable goal into a
 511 reachable one — refutation is preserved. Mechanized as `tolerance_sound`, axiom-free. $\square$

512 **Theorem 3** (Capability monotonicity). *If $\Gamma \subseteq \Gamma'$ and Γ reaches φ_{goal} , then Γ' reaches φ_{goal} .*
 513 *Granting tools never makes an achievable goal impossible.*

514 *Proof.* Enlarging the capability context only *adds* guarded effects: since $\Gamma \subseteq \Gamma'$, every capability
 515 firing available under Γ is available under Γ' with the same guard and effect (WORLD-ACT in B.1
 516 quantifies over $a \in \text{dom}(\Gamma) \subseteq \text{dom}(\Gamma')$), so every step of the Γ -transition relation is a step of
 517 the Γ' -relation. Suppose Γ reaches the goal configuration C from C_0 . By Lemma 2, the same
 518 configuration is reachable under Γ' , and $\text{goal}_{\varphi_{\text{goal}}}(C)$ is unchanged because the goal predicate does
 519 not depend on Γ . Hence Γ' reaches φ_{goal} via the same witnessing run. The other premises of ACH
 520 are monotone in the same direction ($\text{caps}(G) \subseteq \text{dom}(\Gamma) \subseteq \text{dom}(\Gamma')$ still holds, and realizability
 521 and conformance do not mention Γ beyond $a \in \text{dom}(\Gamma)$), so the whole judgment is preserved. The
 522 transport half is mechanized as `cap_monotone`, axiom-free, stated over an inclusion of step relations;
 523 the reduction of $\Gamma \subseteq \Gamma'$ to that inclusion is the WORLD-ACT argument above, on paper. $\square$

524 **Concrete instance.** The flight skill of Section 2, variant A, is mechanized as `FlightInstance`. We
 525 prove `flight_refuted` (no run reaches `booked` $\wedge$ `confirmation_sent`, since no step touches `conf`)
 526 and `booking_reachable` (a booking is still reachable). Both are axiom-free, witnessing non-vacuity
 527 of Theorem 1.

528 D.3 Operational correspondence: subject reduction and session fidelity

529 The direct judgment $G \vdash (\mathbb{M}; W)$ is not merely preserved by reduction: the session and its protocol
 530 step in *lockstep* over the labelled transition systems of Section B.2. This is the standard soundness
 531 backbone of a session calculus [13, 11], obtained here *directly*, with no projection and no merge. For
 532 world-changing interleavings, the statements below assume two explicit frame conditions: effects of
 533 participant-disjoint actions commute, and an action moved under a pending marker preserves that
 534 marker's formula. Communication-only reductions satisfy both conditions vacuously.

535 **Lemma 3** (Residual capability). *If $(G, W) \rightsquigarrow_{\Lambda} (G', W')$, then $\Lambda \in \text{cap}(G)$.*

536 **Lemma 4** (Participant agreement). *If $G \vdash (\mathbb{M}; W)$, then $\text{prt}(G) = \text{prt}(\mathbb{M})$.*

537 Both follow by induction on the transition or typing derivation, respectively; they are the capability-
 538 membership and participant invariants used in the interleaving cases below.

539 **Theorem 4** (Subject reduction). *If $G \vdash (\mathbb{M}; W)$ and $(\mathbb{M}, W) \xrightarrow{\Lambda} (\mathbb{M}', W')$, then $(G, W) \rightsquigarrow_{\Lambda}$*
 540 *(G', W') for some G' with $G' \vdash (\mathbb{M}'; W')$.*

541 **Theorem 5** (Session fidelity). *If $G \vdash (\mathbb{M}; W)$ and $(G, W) \rightsquigarrow_{\Lambda} (G', W')$, then $(\mathbb{M}, W) \xrightarrow{\Lambda}$*
 542 *$(\mathbb{M}', W')$ for some $\mathbb{M}'$ **and** W' with $G' \vdash (\mathbb{M}'; W')$.*

543 We prove both by induction on the typing derivation $G \vdash (\mathbb{M}; W)$, with a case analysis on the
 544 transition; the one auxiliary fact used throughout is that typing respects pointwise-equal sessions
 545 (`ctypes_ext`), which lets the interleaving cases reassociate independent updates.

546 *Proof of Theorem 4 (subject reduction).* By induction on $G \vdash (\mathbb{M}; W)$.

547 *Case T-END* ($G = \text{End}$, every role idle). Then $\mathbb{M}$ has no output or pending action, so the only
 548 available transition is the goal observation S-GOAL at $\Lambda = \check{\varphi}_{goal}$ when $W \models \varphi_{goal}$, which leaves
 549 $(\mathbb{M}, W)$ fixed; there is no residual type to preserve and the claim holds (otherwise no Λ -transition
 550 exists and it holds vacuously).

551 *Case T-GOAL* ($G = \check{\varphi}.G_0$, with $W \models \varphi$ and $G_0 \vdash (\mathbb{M}; W)$, $\text{prt}(G_0) = \text{prt}(\mathbb{M})$). The transition
 552 $(\mathbb{M}, W) \rightarrow_{\Lambda} (\mathbb{M}', W')$ splits on its label.

- 553 • *Goal observation* ($\Lambda = \check{\varphi}$, by S-GOAL, so $\mathbb{M}' = \mathbb{M}$ and $W' = W$). Since $W \models \varphi$, rule
 554 G-GOAL-E discharges the marker, $(\check{\varphi}.G_0, W) \rightsquigarrow_{\check{\varphi}} (G_0, W)$, and $G_0 \vdash (\mathbb{M}; W)$ is the
 555 premise: take $G' = G_0$.
- 556 • *Interleaving* ($\Lambda \neq \check{\varphi}$). A goal marker is not a session construct, so this is already a
 557 transition of G_0 's session; the induction hypothesis gives $(G_0, W) \rightsquigarrow_{\Lambda} (G', W')$ with
 558 $\Lambda \in \text{cap}(G_0)$ and $G' \vdash (\mathbb{M}'; W')$. When the move is world-preserving (a communica-
 559 tion, $W' = W$), G-GOAL-I carries the still-pending marker along, $(\check{\varphi}.G_0, W) \rightsquigarrow_{\Lambda}$
 560 $(\check{\varphi}.G', W)$, and $\check{\varphi}.G' \vdash (\mathbb{M}'; W)$ by T-GOAL (the side condition $\text{prt}(G') = \text{prt}(\mathbb{M}')$ is
 561 preserved because Λ acts on the same participants on both sides). When the move is *world-*
 562 *changing* (a capability firing, $W' \neq W$), G-GOAL-I still commutes it under the marker,
 563 $(\check{\varphi}.G_0, W) \rightsquigarrow_{\Lambda} (\check{\varphi}.G', W')$; re-typing the residual by T-GOAL then additionally asks
 564 $W' \models \varphi$. This follows from the goal-stability hypothesis on actions commuting under pend-
 565 ing markers. The global type sequences the marker before G_0 's actions whereas the session
 566 may fire the action first; the interleaving rule lets the type catch up under that hypothesis.
 567 Removing the hypothesis, rather than this conditional case, is the open correspondence
 568 problem recorded in Section 7.

569 *Case T-COMM*, with head $p \rightarrow q$ and branches G_i . Here $\mathbb{M}$ is $p[q! \dots] \parallel q[p? \dots] \parallel \mathbb{M}_0$, and each
 570 branch satisfies $G_i \vdash (p[P_i] \parallel q[Q_i] \parallel \mathbb{M}_0; W)$. If the transition is a goal observation $\Lambda = \check{\varphi}$
 571 (S-GOAL, leaving $(\mathbb{M}, W)$ fixed), it is matched by G-COMM-I: $\text{prt}(\check{\varphi}) = \emptyset$ makes the disjointness
 572 premise trivial, and each branch realises the same observation by the induction hypothesis, using
 573 $\check{\varphi} \in \text{cap}(G_i)$ — the observation concerns a marker the protocol actually pends (see the note below).
 574 Otherwise the transition is a communication $\Lambda = r \ell t$. Because p offers only outputs to q and q only
 575 inputs from p , the sender/receiver shapes force a dichotomy.

- 576 • *Head move* ($r = p$). Then $t = q$ and $\ell = \ell_k \in I$, and $\mathbb{M}' = p[P_k] \parallel q[Q_k] \parallel \mathbb{M}_0$. Rule
 577 G-COMM-E gives $(G, W) \rightsquigarrow_{p\ell_k q} (G_k, W)$, and $G_k \vdash (\mathbb{M}'; W)$ is exactly the k -th premise.
- 578 • *Interleaving move* ($r \neq p$). The shapes then force $r, t \notin \{p, q\}$: r is an output so $r \neq q$,
 579 and t is an input so $t \neq p$, while $t = q$ would force $r = p$. Since r, t lie in the bystanders
 580 $\mathbb{M}_0$, the *same* communication is enabled in each branch's configuration $p[P_i] \parallel q[Q_i] \parallel \mathbb{M}_0$
 581 (the head updates do not touch r, t). The induction hypothesis at each i yields G'_i with
 582 $(G_i, W) \rightsquigarrow_{\Lambda} (G'_i, W)$ and $G'_i \vdash (p[P_i] \parallel q[Q_i] \parallel \mathbb{M}'_0; W)$, where $\mathbb{M}'_0$ is $\mathbb{M}_0$ after the
 583 bystander step. As $\text{prt}(\Lambda) = \{r, t\}$ is disjoint from $\{p, q\}$, rule G-COMM-I assembles
 584 $(G, W) \rightsquigarrow_{\Lambda} (p \rightarrow q: \{\ell_i, G'_i\}, W)$, and re-applying T-COMM (the sender/receiver at p, q
 585 are untouched, and each branch is retyped by the induction hypothesis, reassociating the
 586 independent updates via `ctypes_ext`) gives the residual typing.

587 Communications leave the world fixed, so no independence hypothesis is needed. For a world-
 588 changing action at the head (T-ACT), G-ACT-E matches and the residual G is typed at the *same*

post-effect world the rule reaches — this is exactly where the lockstep world pairing of the corrected T-ACT is used; a goal observation at an action-typed configuration is matched by G-ACT-I exactly as in the T-COMM case, since $\text{prt}(\checkmark\varphi) \cap \{p\} = \emptyset$ trivially. Interleaving *two* actions requires their effects to commute ($\text{eff}_a \circ \text{eff}_b = \text{eff}_b \circ \text{eff}_a$ for disjoint a, b), a frame-independence side condition that participant-disjointness alone does not supply; under it the T-ACT interleaving case goes through identically. $\square$

Proof of Theorem 5 (session fidelity). By induction on $G \vdash (\mathbb{M}; W)$, inverting the global step $(G, W) \rightsquigarrow_{\Lambda} (G', W')$. T-END admits only the goal observation G-GOAL-E at $\checkmark\varphi_{\text{goal}}$, matched on the session by S-GOAL. For T-GOAL, the step is either G-GOAL-E — matched by S-GOAL on the session, the marker discharged — or G-GOAL-I, whose premise $(G_0, W) \rightsquigarrow_{\Lambda} (G'_0, W')$ the induction hypothesis realises as an underlying session step that carries the marker along. For T-COMM, the global step is either G-COMM-E — and the prescribed head communication $p \ell_k q$ is enabled by S-COMM, landing in the k -th premise — or G-COMM-I with $\text{prt}(\Lambda) \cap \{p, q\} = \emptyset$; then the induction hypothesis on any branch produces a bystander communication, whose endpoints lie outside $\{p, q\}$, so S-COMM fires it on $\mathbb{M}$ and T-COMM retypes the result. The action cases mirror those of Theorem 4. $\square$

Mechanization. Both theorems are mechanized axiom-free for the single-label communication fragment — the fragment in which bystander interleaving is unconditional — in `DirectTypingSR.v`, over labelled session and global transition systems, with no projection, merge, or local type:

```

608   Theorem subject_reduction : forall W G s s' L,
609     ctypes W G s -> lstep L s s' ->
610     exists G', gstep W L G G' /\ ctypes W G' s'.
611   Theorem session_fidelity : forall W G G' s L,
612     ctypes W G s -> gstep W L G G' ->
613     exists s', lstep L s s' /\ ctypes W G' s'.
614   (* Print Assumptions => Closed under the global context (both). *)

```

The head-move case for world-changing *actions* is additionally mechanized in `DirectTyping.v` (`type_directed_safety`, also `progress`). Observable goal labels appear in neither Coq development: the mechanized label alphabet is communication-only, and its goal rule fuses marker discharge with a continuation step. Participant-matching side conditions, labelled action interleaving, and multi-branch choice are likewise absent. Completing a faithful mechanization of the paper rules remains work in Section 7.

Concrete instance. `HandoffInstance` mechanizes the planner/worker handoff of Section 1 on both sides. The *good* pairing — planner sends, worker delivers, goal checked — is proved `ctypes`-typed (`good_typed`) and its run is proved to reach a world satisfying the goal (`good_runs_to_goal`). The *bad* pairing — both roles start with an input, exactly the deadlock of Section 1 — is proved `Stuck` (`bad_stuck`), and the contrapositive of `progress` then gives, for free, that *no non-trivial global type can ever type it* (`bad_untimeable`): the rule rejects the classical deadlocking handoff without any projection ever being computed.

628 D.4 Incompleteness, by design

Proposition 1 (Incompleteness). $\Gamma; G \models \Diamond_{\text{init}} \varphi_{\text{goal}}$ may hold while no concrete runtime execution reaches φ_{goal} : the witnessing abstract path can be spurious when α collapses a distinction that mattered. The system is sound for $\times$ and incomplete for $\checkmark$.

This is the price of decidable, tolerant over-approximation; tightening α trades tolerance for precision. The residue is confined to two edges, motivating the layered architecture:

- 634 • **Payload faithfulness** (bottom edge): a tool’s declared post-condition (“filters to fares
635 < 500”) may be false — a runtime obligation for a deterministic monitor (Layer C), not a
636 static one.
- 637 • **Intent fidelity** (top edge): the formalized goal may not capture what the user meant —
638 irreducibly semantic, surfaced at the single compaction checkpoint for human review.

639 D.5 Decidability and the autonomy boundary

640 **Theorem 6** (Decidability). *In the fragment with finitely many roles (static topology), regular global*
 641 *types represented by finite tail-recursive control graphs, decidable state logic $\mathcal{L}$, and $\mathcal{L}$ -definable*
 642 *effects, $\Gamma; G \models \Diamond_{init} \varphi_{goal}$ is decidable.*

643 *Proof.* The pack-level achievability judgment $\Gamma; G \models \Diamond_{init} \varphi_{goal}$ (Section B.2) is an existential
 644 reachability query over configurations (G', W) ; we exhibit a terminating decision procedure and
 645 argue its completeness.

646 *The search space is finite.* A configuration has two components. (i) *Control*, the residual global
 647 type G' reachable from G by $\rightsquigarrow$. Since G is a regular tree (tail-recursion: finitely many distinct
 648 subtrees, Section A.3) and each $\rightsquigarrow$ step either descends into a subtree (G-COMM-E, G-ACT-E) or
 649 erases a marker (G-GOAL-E), the set $\{G' : (G, _) \rightsquigarrow^* (G', _)\}$ is a subset of the finitely many
 650 subtrees of G , hence finite. (ii) *Symbolic world*. The boolean part is a valuation $B : \text{Pred} \rightarrow \mathbb{B}$
 651 over the finitely many predicates named in the pack, so there are $2^{|\text{Pred}|}$ boolean states. The numeric
 652 part is summarized by an accumulated constraint $\psi \in \mathcal{L}$ (a conjunction of the guards taken and
 653 the Nd -postconditions assumed); on a control-graph back edge the reference procedure applies the
 654 widening $\psi \mapsto \top$, forgetting accumulated numeric constraints. This one-element numeric summary
 655 is an over-approximation and, together with the finite boolean valuations, leaves only finitely many
 656 symbolic worlds. The product with finite control is therefore finite.

657 *Each edge is computable.* Firing a capability requires checking $W \models pre_a$ and forming sp ; testing
 658 goal-satisfaction requires $W \models \varphi$. Both are entailment/satisfiability queries in quantifier-free linear
 659 integer arithmetic with uninterpreted predicates — a decidable theory, discharged by an SMT solver.
 660 Branch selection (REACH-OR) is a finite disjunction. Existence of an initial world is decided by the
 661 same SMT query over $init$; each satisfying symbolic initial valuation seeds the finite search.

662 *Termination and correctness.* Forward search (breadth-first over $\rightsquigarrow$) maintains a visited set of (control,
 663 symbolic-world) pairs; because that set is finite it adds each pair at most once, so the search halts. It
 664 answers ACHIEVABLE iff it enumerates a goal-satisfying configuration, and IMPOSSIBLE otherwise.
 665 Soundness of the IMPOSSIBLE verdict is Theorem 1; and widening only *enlarges* the reachable set (it
 666 weakens constraints, adding edges), so by Theorem 2 it never introduces a false refutation. Hence the
 667 procedure is a decision procedure for $\Gamma; G \models \Diamond_{init} \varphi_{goal}$. $\square$

668 **Proposition 2** (Undecidability in a dynamic-topology extension). *Consider the extension G^+ of the*
 669 *grammar with spawn $r.G$, which creates fresh participants at run time. With an unbounded role*
 670 *population and dynamic channels, achievability in G^+ is undecidable.*

671 *Proof sketch.* By reduction from the control-state reachability problem for communicating finite-
 672 state machines (CFSMs), which is undecidable [15]. Fix a CFSM instance: finite-control machines
 673 M_1, M_2 exchanging messages over unbounded FIFO channels, and a target control state q_f of M_1 ;
 674 deciding whether a run reaches q_f is undecidable.

675 We compute, from this instance, a skill whose goal is achievable iff the CFSM reaches q_f . Each
 676 machine M_i becomes a role whose process mirrors its control graph: a transition that sends message
 677 m spawns, at run time, a fresh *courier* participant carrying m (this is exactly the dynamic-spawning
 678 capability the theorem hypothesizes), and a transition that receives m synchronizes with the oldest
 679 outstanding courier for that channel, so the unbounded family of couriers realizes the unbounded FIFO.
 680 The required order can be made explicit by assigning sequence numbers through the numeric world
 681 state and guarding receives on the next sequence number; the *Asg/Nd* machinery of Section A.2
 682 supplies those updates. A boolean predicate at_{q_f} is established by the single capability guarding
 683 M_1 's entry into q_f , and we set $\varphi_{goal} = at_{q_f}$. By construction a run of the skill reaches a state
 684 satisfying φ_{goal} iff the CFSM has a run reaching q_f : the courier discipline then puts the skill's
 685 reachable configurations in correspondence with the CFSM's channel contents and control states.
 686 The construction is a computable map from CFSM instances to packs, so a decision procedure for
 687 achievability would decide CFSM control-state reachability — impossible. Hence achievability is
 688 undecidable once participants may be spawned dynamically. (The finiteness argument of Theorem 6
 689 breaks at exactly the point this reduction exploits: an unbounded number of couriers makes the set of
 690 reachable controls infinite.) $\square$

The two results isolate one concrete boundary: static finite topology admits a total decision procedure, while unbounded dynamic spawning does not. This is a boundary on topology, not a claim that all forms of agent autonomy are undecidable. The implementation returns UNKNOWN when a pack violates the static-topology side condition of Theorem 6, unless an independent capability or establisher refutation has already proved it impossible.

D.6 The partition of obligations

Corollary 1. *The architecture separates “will the goal be achieved” into four obligations on disjoint substrates:*

Concern	Where	Mechanism
goal achievability	static, decidable	Coq specification; deterministic search; no LLM
per-step safety / least privilege	Layer B (future)	same engine, $\square$ + permission tiers
payload faithfulness	Layer C (future), run-time	deterministic monitors with blame
intent fidelity	top edge, synthesis	one human checkpoint

D.7 Verifiable feedback beyond a scalar reward

The verdict is a heterogeneous verification signal rather than a scalar score. An IMPOSSIBLE result identifies which premise failed and carries a blocking frontier; ACHIEVABLE carries a witness path and explicitly retains the two deferred obligations; UNKNOWN is an abstention outside the decidable fragment, not a refutation. These structured outcomes can serve as cheap negative labels for automated skill or prompt search without an LLM judge, while the intent-fidelity checkpoint keeps the irreducible semantic decision with a human. We treat such search as an application of the verifier, not as a contribution evaluated here.

E Token economics of pre-execution refutation

The practical case for deciding achievability *before* execution is an economic one, so we state it in tokens rather than adjectives. Two asymmetries carry it. First, no model sits in the trusted path: the schema gate, projection, conformance and Z3 reachability spend **zero tokens**, and the deterministic front-end spends zero as well. Only the optional LLM compaction spends tokens, and it is paid *once per skill version*. Second, an agent turn re-sends its whole context: writing S for the harness prompt, K for the loaded skill, g for the mean tokens appended per turn and o for output per turn, a run of T turns reads $T(S+K) + gT(T-1)/2$ input and To output tokens. Compaction is linear in the skill and paid once; a doomed run is *quadratic* in its turns and recurs on every invocation.

Runtime waste is necessarily a model — it prices a run that, if the refutation is correct, never happens — so we publish the model rather than a single number. Each refutation reason carries a (low, typical, high) band for how long an agent flails before that failure stops it, and how many agents are billed. With conservative defaults ($S=2500$, $g=700$, $o=250$) and a median real consumer skill ($K \approx 1.7K$ tokens), compaction costs $\sim 2.8K$ tokens against the following avoided waste per prevented invocation:

Verdict reason	turns (lo–typ–hi)	waste/run	leverage
MISSING_CAPABILITY	3–8–14, 1 agent	50 240	18×
GOAL_UNSAT	8–18–30, 1 agent	176 040	63×
NON_CONFORMANT	6–15–30, 2 agents	261 900	94×
NON_PROJECTABLE	6–15–30, 2 agents	261 900	94×
BLOCKED_GUARD	12–25–50, 1 agent	305 750	110×

The ordering is the robust part, and it is the expected one. MISSING_CAPABILITY is cheapest at run time: the agent calls a tool that is not there and learns immediately, and most of the cost is the improvisation that follows. BLOCKED_GUARD is dearest, because a precondition no run can satisfy is precisely the case an agent cannot learn from — it retries to the harness turn cap. The two multi-party reasons bill two contexts. GOAL_UNSAT additionally carries a cost the token bill

729 does not show: a run whose goal conjunct has no establisher can terminate *believing* it succeeded.
 730 DYNAMIC_TOPOLOGY claims no savings at all, because an abstention prevents nothing.

731 Over the 15-spec corpus the seven structural refutations cost 12.2K tokens of compaction and avoid
 732 $\sim 1.07\text{M}$ tokens per round of invocations (band 0.28M–3.15M): $\sim 87\times$ leverage, or nothing at all
 733 on the deterministic path. Break-even therefore falls *inside* the first prevented run in every mode
 734 — between 0.9% and 5.5% of it. The honest denominator for a healthy skill, where the check buys
 735 nothing and still costs something, is 2.9% of one successful run over the corpus and 4.0% for a
 736 median real skill; with the deterministic front-end, zero. Prompt caching changes the price of the
 737 wasted tokens, not their number. Halving every waste default leaves leverage at $9\times$ – $55\times$ and the
 738 ordering unchanged. The model, its defaults and their justification are in the artifact (`skillc.tokens`,
 739 `skillc.cost`); the Coq development costs no tokens and is amortized over every skill ever checked.

740 NeurIPS Paper Checklist

741 1. Claims

742 Question: Do the main claims made in the abstract and introduction accurately reflect the
 743 paper’s contributions and scope?

744 Answer: [\[Yes\]](#)

745 Justification: The abstract and Section 1 state that refutation is sound relative to the declared
 746 pack, distinguish UNKNOWN from refutation, and identify the mechanized and paper-only
 747 proof fragments.

748 2. Limitations

749 Question: Does the paper discuss the limitations of the work performed by the authors?

750 Answer: [\[Yes\]](#)

751 Justification: The Limitations and Future Work section discusses relative soundness, static
 752 topology, abstraction coarseness, mechanization gaps, and the proof-of-concept corpus.

753 3. Theory assumptions and proofs

754 Question: For each theoretical result, does the paper provide the full set of assumptions and
 755 a complete (and correct) proof?

756 Answer: [\[Yes\]](#)

757 Justification: Appendix D states the simulation, stability, commutativity, finiteness, and
 758 topology assumptions with proofs or proof sketches; the artifact checks the explicitly
 759 identified Coq fragments.

760 4. Experimental result reproducibility

761 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
 762 perimental results of the paper to the extent that it affects the main claims and/or conclusions
 763 of the paper (regardless of whether the code and data are provided or not)?

764 Answer: [\[Yes\]](#)

765 Justification: Sections 4–5 describe the deterministic checker, verdict semantics, corpus
 766 construction, binary accounting, and separate treatment of abstentions.

767 5. Open access to data and code

768 Question: Does the paper provide open access to the data and code, with sufficient instruc-
 769 tions to faithfully reproduce the main experimental results, as described in supplemental
 770 material?

771 Answer: [\[Yes\]](#)

772 Justification: The anonymized supplemental artifact contains the checker, synthetic corpora,
 773 proof developments, pinned continuous-integration workflow, and reproduction commands.

774 6. Experimental setting/details

775 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
 776 rameters, how they were chosen, type of optimizer) necessary to understand the results?

777 Answer: [\[N/A\]](#)

778 Justification: The evaluation is a deterministic execution over declared packs; it performs no
779 training, optimization, sampling, or data splitting.

780 **7. Experiment statistical significance**

781 Question: Does the paper report error bars suitably and correctly defined or other appropriate
782 information about the statistical significance of the experiments?

783 Answer: [N/A]

784 Justification: Each corpus case and checker result is deterministic, so repeated trials have no
785 sampling variance; the paper reports exact counts rather than inferential statistics.

786 **8. Experiments compute resources**

787 Question: For each experiment, does the paper provide sufficient information on the com-
788 puter resources (type of compute workers, memory, time of execution) needed to reproduce
789 the experiments?

790 Answer: [Yes]

791 Justification: Section 5 states that evaluation is CPU-only, requires no model inference or
792 training, and gives the corpus size and solver timeout governing the run.

793 **9. Code of ethics**

794 Question: Does the research conducted in the paper conform, in every respect, with the
795 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

796 Answer: [Yes]

797 Justification: The work uses synthetic packs, no personal data or human subjects, and makes
798 bounded claims with explicit trust assumptions and limitations.

799 **10. Broader impacts**

800 Question: Does the paper discuss both potential positive societal impacts and negative
801 societal impacts of the work performed?

802 Answer: [Yes]

803 Justification: The Broader impact paragraph discusses reduced wasted execution as well as
804 false confidence from inaccurate packs and the human/runtime mitigations required.

805 **11. Safeguards**

806 Question: Does the paper describe safeguards that have been put in place for responsible
807 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
808 image generators, or scraped datasets)?

809 Answer: [N/A]

810 Justification: The artifact releases no trained model, scraped dataset, or other high-risk asset;
811 it contains a deterministic checker and synthetic examples.

812 **12. Licenses for existing assets**

813 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
814 the paper, properly credited and are the license and terms of use explicitly mentioned and
815 properly respected?

816 Answer: [N/A]

817 Justification: The evaluation does not reuse an external dataset or model; prior methods and
818 software foundations are credited through citations.

819 **13. New assets**

820 Question: Are new assets introduced in the paper well documented and is the documentation
821 provided alongside the assets?

822 Answer: [Yes]

823 Justification: The supplemental artifact documents the checker interface, pack schema,
824 synthetic corpora, proof scope, limitations, and reproduction workflow.

825 **14. Crowdsourcing and research with human subjects**

826 Question: For crowdsourcing experiments and research with human subjects, does the paper
827 include the full text of instructions given to participants and screenshots, if applicable, as
828 well as details about compensation (if any)?

829 Answer: [N/A]

830 Justification: The research involves neither crowdsourcing nor human-subject experiments.

831 **15. Institutional review board (IRB) approvals or equivalent for research with human**
832 **subjects**

833 Question: Does the paper describe potential risks incurred by study participants, whether
834 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
835 approvals (or an equivalent approval/review based on the requirements of your country or
836 institution) were obtained?

837 Answer: [N/A]

838 Justification: The research involves no study participants or human-subject data.

839 **16. Declaration of LLM usage**

840 Question: Does the paper describe the usage of LLMs if it is an important, original, or
841 non-standard component of the core methods in this research? Note that if the LLM is used
842 only for writing, editing, or formatting purposes and does not impact the core methodology,
843 scientific rigor, or originality of the research, declaration is not required.

844 Answer: [N/A]

845 Justification: The paper’s core method does not depend on an LLM. The main contributions—
846 the achievability type discipline, its formal metatheory, and the deterministic checker—are
847 defined and verified independently of LLM inference. The LLM is used only as a front-end
848 to compact a natural-language skill into the formal pack consumed by the checker; the
849 same formal pack can be provided directly without an LLM. The formal guarantees and the
850 reported evaluation therefore do not depend on the behavior of any particular LLM.