

Affordable Mistakes: Severity-Aware Multiparty Session Types for Participants that Choose Wrongly*

Anonymous author

Anonymous affiliation

Anonymous author

Anonymous affiliation

Abstract

Multiparty session types guarantee that well-typed participants interact correctly, on the hypothesis that each participant *is* its type. A participant that was instructed rather than compiled — a language-model agent — breaks that hypothesis in one specific way: at a branch where it should select *A* it may select *B*. Forbidding this is impossible, and would be useless if it were possible. We change what the discipline is for: it does not prevent wrong choices, it prevents catastrophic ones and reports the rest as risk.

The fault model is the participant's own selection at a branch, defined through the guards that asserted global types already carry. Cascading is a budget *k* of wrong selections, with no probabilities anywhere. Budgeted deviation is itself old — reactive synthesis and robustification have owned it for a decade. What is new is carrying a quantity a participant *spends* across a reduction relation and into typed programs: a session typed against a *k*-tolerant protocol by the base discipline's conformance judgment is hazard-free on every run of cost at most *k*, budgets distribute over participants, and the guarantee survives recursion and every permutation of independent nodes. One syntax-directed rule characterizes the condition exactly; each wrong selection is classified **Benign**, **Futile** or **Catastrophic** by outcome against a STRIPS-style world, and these are three consecutive intervals of budgets, not an arbitrary carve-up. The condition composes against an interface rather than a protocol's internals, and is decidable for regular protocols. Everything is mechanized in Coq, axiom-free, and the tool's decision on the boolean fragment is a kernel extracted from that proof.

The evaluation asks whether any of this matters to a real agent. On a benchmark of seventeen protocols the analysis names the point-of-no-return action of every 0-tolerant one; 340 runs of two production models find every observed catastrophe consistent with the theorem and the catastrophe rate ordered by tolerance degree; and on 162 skills from thirteen public repositories, run in two runtimes with every artifact verified by recomputation, the refusals coincide with the runs that ended in wasted tokens or a fabricated result — at realistic input sizes no refuted run produced a correct artifact.

2012 ACM Subject Classification Theory of computation → Type structures; Theory of computation → Process calculi; Software and its engineering → Software verification

Keywords and phrases Session types, Multiparty, Behavioural types, Safety, Irreversibility, Fault tolerance, Mechanized metatheory

Digital Object Identifier 10.4230/LIPIcs.ECOOP..0

Supplementary Material [Anonymous supplementary material](#)

Funding [Anonymous funding](#)

Acknowledgements [Anonymous acknowledgements](#)

* **WORKING DRAFT** — not for submission. Section 3.



© [Anonymous author\(s\)](#);

licensed under Creative Commons License CC-BY 4.0

Leibniz International Proceedings in Informatics

LIPICS Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

WIP: Status. Everything in Sections 5–9 is mechanized (proof/, 100 results — 86 for the present theory and 14 for the audit of §14 — all **Print Assumptions** closed). The evaluation numbers in Section 10 are real (scripts/severity_eval.py, results in paper/WIP/results/). Scope caveats are stated where they apply. The predecessor draft was withdrawn after its own mechanized audit found it vacuous (Section 14).

1 Introduction

Multiparty session types (MPST) guarantee communication safety, session fidelity and deadlock-freedom for systems whose participants *are* their local types [1]: the code was checked, so every run is a compliant run. A participant that was instructed rather than compiled breaks that hypothesis in a specific way. At a choice point it may take branch *B* where the situation demanded branch *A* — not by violating the protocol (both branches are in the type) but by judging the situation wrongly. This is the ordinary condition of a language-model agent occupying a role, and it is why such agents are employed: their latitude.

Nothing here is specific to language models. The hypothesis MPST rests on — the participant *is* its type — fails for every participant that was instructed rather than compiled: an operator following a runbook at a migration console, a third-party service behind an adapter whose behaviour is a contract rather than a checked artifact, a learned controller. Each occupies a role and may, at a branch the protocol offers, take the one the situation forbids. Language-model agents are the case that makes the question urgent and the case our evaluation can drive at scale, but the discipline asks only that the participant choose, not how.

The tempting response is to make the type system prevent it. That is doubly wrong: impossible, because the judgement is the participant's, not the protocol's; and undesirable, because a participant that may never choose wrongly is one that may never choose. What is needed is *triage*.

The reframing. Do not ask whether the participant complied. Ask *what a wrong choice costs*. Three answers matter, and the middle one is what existing disciplines cannot express (Figure 1):

- Benign — a detour; the goal is still reachable.
- Futile — the goal is lost, but nothing is harmed. *Failure is not disaster*.
- Catastrophic — a hazard becomes reachable.

A system that blocks Futile is useless; one that permits Catastrophic is dangerous. Separating them is the whole job. Thesis: *session types say what may happen; this says what you can survive*.

What is and is not new. The ingredients are old; the discipline is not. Branch guards on global types are design-by-contract MPST [2, 4, 5], and guard violation on selection is the alarm condition of the monitoring line [3]. Whether a goal remains reachable is dead-end detection in planning [28, 29]; whether an effect can be undone is undoability checking [30, 31]. Bounded fault tolerance is classical in planning [26, 27] and in MPST with crash-stop failures [8, 9]. A syntax-directed system equivalent to a model checker is a known template [20, 21]. We claim none of these. What we claim is C1–C5 below, and nothing wider.

We are careful with the word *type system*. The condition is on global types, as well-assertedness [2] and deadlock-freedom conditions are; the participants are typed by the base discipline’s direct-conformance judgment, unchanged. The type-theoretic content is the theorem that couples the two (Section 7), not the rule that decides the condition.

Contributions. Throughout, ✓*name* marks a result proved in the accompanying Coq development and names its identifier there.

- C1** *A budget a typed participant spends, carried to programs* (§4, §7). A session typed against a k -tolerant protocol is hazard-free on every run whose misselection cost is at most k (✓*bridge_run*); the budget distributes over participants, so a global allowance is a contract that can be split and checked per role (✓*budget_distributes*); the guarantee survives recursion (✓*bridge_mu*) and every permutation of independent nodes, with the side conditions discharged syntactically (✓*bridge_interleaved*, ✓*strips_swap_act*). Budgeted deviation is not ours (§13); *carrying a quantity a participant spends across a reduction relation, into typed programs*, is what no prior discipline does.
- C2** *The condition that makes C1 checkable* (§6): one syntax-directed rule, sound and complete for k -tolerance (✓*TC_exact*), which is exactly non-reachability in a product graph (✓*TC_regular*).
- C3** *An ordered diagnosis* (§5): every wrong choice is Benign, Futile or Catastrophic, and these are not an arbitrary carve-up — the class of a residual is monotone in the budget along the chain Futile < Benign < Catastrophic (✓*severity_monotone_in_budget*), so the three classes are consecutive intervals of budgets and k^* is a genuine threshold (✓*tolerance_degree_is_a_threshold*). Budgeted reachability is exactly “some run with at most k wrong choices reaches a hazard” (✓*reach_mu_iff_run*), so the classes are statements about executions.
- C4** *Decidability and modularity* (§8, §9): the μ -unfolding closure is finite (✓*cands_closed*) and the decision procedure is proved sound and complete (✓*decide_mu_correct*); sequential composition is typed against an interface that carries the remaining budget, not the internals (✓*TC_seq_interface*). Four repairs are proved sound, two exactly.
- C5** *Evidence* (§10): an implementation whose verdict agrees with a kernel extracted from the proof on 499 of 499 random protocols; a severity benchmark; 340 live-agent runs showing the Catastrophic verdicts name branches real agents take; and 134 verified runs over 162 real skills in two runtimes showing the achievability verdicts predict which runs are wasted or fabricated, with the false-refutation rate measured.


A design expectation was refuted by the mechanization and a benchmark finding refines it: tolerance is anti-monotone in the capability context for a fixed hazard (✓*tolerance_antitone_in_ctx*), yet a new tool that re-establishes a lost atom can *remove* a point of no return (§10).

2 Overview

Running example. A planner *p* and a worker *q* holding a purchase tool. Atoms *verified*, *booked*, *funds*; *verify* adds the first, *purchase* adds the second and deletes the third, and no tool restores *funds*. The protocol offers two branches:

$$G_{\text{bad}} = p \rightarrow q : \{ \text{safe. verify@q. purchase@q. End} , \text{fast. purchase@q. End} \}$$

with goal $\varphi = \text{booked} \wedge \text{verified}$. Under the default instantiation (§4) *safe* is intended because the goal survives it and *fast* is a misselection because it does not. The tool reports, and



■ **Figure 1** The idea, on the running example of §2. At a guarded choice in world W the branch whose guard holds is intended; taking another is a misselection. Its severity is a property of the residual, and the three fates drawn are the only three: the goal still reachable (**Benign**), the goal lost with no hazard reachable (**Futile**), or a hazard reachable (**Catastrophic**) — here by crossing a point of no return, an irreversible capability after which the goal cannot be recovered. In G_{bad} the misselected branch takes the third.


the Coq development proves for the same instance: G_{bad} is 0-tolerant ($\checkmark G_{\text{bad_is_0_tolerant}}$) — if the participant never misselects, nothing bad happens, which is where classical MPST stops; it is *not* 1-tolerant ($\checkmark G_{\text{bad_not_1_tolerant}}$) — one wrong choice purchases unverified, and **purchase** is the point of no return; and narrowing away **fast** is tolerant at every k ($\checkmark G_{\text{good_is_k_tolerant}}$). The headline number is the *tolerance degree* k^* , the largest k the protocol survives.

For a protocol with two decision points the tool’s output is what an engineer needs (Section 10, `migration_backup`):

```
migration_backup: tolerance degree k* = 1
choice nodes 4, branches 5, irreversible ['drop_old']
/choice@ops#0
  backup          Benign      intended
  skip_backup     Catastrophic misselection PNR=drop_old
/choice@ops#0/backup/choice@ops#2
  drop_old        Benign      intended
  keep_old        Benign      intended
/choice@ops#0/skip_backup/choice@ops#1
  drop_old        Catastrophic misselection PNR=drop_old
first catastrophe within budget 2:
  choose/skip_backup > act/migrate > choose/drop_old > act/drop_old
repair (narrow):
  remove skip_backup at /choice@ops#0
  remove drop_old at /choice@ops#0/skip_backup/choice@ops#1
```

“Tolerant of one wrong choice; two wrong choices reach data loss, here is the path and the action that burns the bridge” is actionable. “Impossible” is not.

Architecture. Figure 2: a compile-time check on the pack, before any participant runs, with no language model in the trusted core — the base discipline’s trust boundary. The runtime gate is inherited unchanged: it refuses *out-of-set* attempts (labels the protocol does not offer) before delivery, so those change no state. The deviation this paper analyses is *in-set*: a branch the type permits and the situation forbids.



■ **Figure 2** Architecture. Compaction from prose to a pack is untrusted and may use a language model; nothing beyond the trust boundary does. The severity classifier runs the discipline’s existing reachability search pointwise at every branch of every choice, at every abstract world and budget, and emits a per-branch severity, the tolerance degree k^* , a witness naming the point-of-no-return action, and repair suggestions. The certified pack drives the gated runtime that refuses out-of-protocol attempts.

3 The Base Discipline

We build on a goal-marked, capability-guarded synchronous multiparty discipline, stated here in the detail this paper uses. Predicates $p \in \text{Pred}$, numeric variables $x \in \text{Var}$, roles p, q , capabilities $a \in \text{Cap}$, labels $\ell \in \text{Lab}$. The state logic is quantifier-free linear integer arithmetic with uninterpreted booleans; a world $W = \langle B, N \rangle$ values it, and the checker’s decidable fragment abstracts predicates concretely and numerics symbolically, widening at loop heads so the reachable world set is finite. A capability context maps each possessed tool to a guarded STRIPS effect $\Gamma(a) = \langle pre_a, eff_a \rangle$, where eff_a adds and deletes predicates and assigns numerics. Firing is WORLD-ACT: $\langle W \rangle a \langle W' \rangle$ iff $W \models pre_a$ and $W' \in \text{apply}(eff_a, W)$. We write φ for the goal and $\mathcal{H}$ for the hazard, both predicates over worlds. Global types are coinductive regular trees

$$G ::=^{coind} p \rightarrow q : \{ \ell_i. G_i \}_{i \in I} \mid a @ p. G \mid \checkmark \varphi. G \mid \text{End}$$

and sessions $\mathbb{M} = p_1[P_1] \parallel \dots$ of processes $P ::=^{coind} \mathbf{0} \mid p! \{ \ell_i. P_i \} \mid p? \{ \ell_i. P_i \} \mid a. P$ are typed *directly* against G by the coinductive judgment $G \vdash (\mathbb{M}; W)$ — no local types, projection, or separate subtyping [16, 17] — whose T-COMM requires the sender to offer exactly the protocol’s labels, the receiver at least them, and every branch to check against the same residual. A *pack* $\langle \Gamma, G, \varphi, \mathcal{H} \rangle$ is the checked object: a capability context, a global type, a goal and a hazard. Achievability is existential may-reachability of a goal-satisfying configuration; the base metatheory (subject reduction, session fidelity, refutation soundness) is mechanized for the finite fragment. This paper adds one annotation to choice, one counter to the semantics, and one well-formedness condition.

4 Misselection and the Budget

► **Definition 1** (Guarded choice). *A choice node carries, per branch, the predicate that makes that branch the intended one: $p \rightarrow q : \{ \ell_i[\psi_i]. G_i \}_{i \in I}$. In world W , branch i is intended if*

$W \models \psi_i$; taking a branch with $W \not\models \psi_i$ is a misselection.

The syntax is that of asserted global types [2], whose well-formedness conditions guarantee that a compliant participant can always pick a satisfied branch. We keep the syntax and drop the guarantee: the guards are not a constraint on an implementation but the *definition of wrong*.

Default instantiation. Guards need not be written. When a branch carries none, the implementation uses the *rational-choice* default: a branch is intended in W iff the goal is still reachable from its residual. Under this default a misselection is a choice that forfeits the goal, and once the goal is forfeit every subsequent choice is a misselection — there is no longer a right one — so the budget counts decisions taken after the point at which the task was lost. Explicit guards override the default where the intended branch is a matter of policy rather than reachability (Section 10 uses both).

► **Definition 2** (Budgeted reachability). $(G, W) \rightsquigarrow_b^{\mathcal{H}}$ holds when a hazard state is reachable from (G, W) using at most b misselections (Coq: `reach_haz`): the head rules of the base semantics, plus

$$\begin{array}{c} \text{RH-COMM-OK} \frac{\ell_i[\psi_i].G_i \in \text{brs} \quad W \models \psi_i \quad (G_i, W) \rightsquigarrow_b^{\mathcal{H}}}{(\text{p} \rightarrow \text{q} : \text{brs}, W) \rightsquigarrow_b^{\mathcal{H}}} \\ \\ \text{RH-COMM-DEV} \frac{\ell_i[\psi_i].G_i \in \text{brs} \quad W \not\models \psi_i \quad (G_i, W) \rightsquigarrow_b^{\mathcal{H}}}{(\text{p} \rightarrow \text{q} : \text{brs}, W) \rightsquigarrow_{b+1}^{\mathcal{H}}} \end{array}$$

Compliant progress is free; a misselection costs one unit. A choice may instead be marked as controlled by the environment, in which case every branch is taken demonically at no cost; the implementation supports this and one benchmark protocol uses it, but the mechanization covers only participant-controlled choice, so no theorem in this paper is stated for it.

► **Definition 3** (k -misselection tolerance). G is k -misselection-tolerant for hazard $\mathcal{H}$ under Γ from W_0 if $\neg(G, W_0) \rightsquigarrow_k^{\mathcal{H}} \cdot^1$ 0-tolerance is classical MPST safety.

The budget is possibilistic: a design parameter of the deployment, not a measurement of the participant. Every theorem below holds for every participant that misselects at most k times, however it does so.

5 Severity and the Point of No Return

Let φ be the goal and $\mathcal{H}$ the hazard; $\rightsquigarrow_b^{\varphi}$ is budgeted goal reachability.

► **Definition 4** (Severity). For a residual (G, W) at budget b : *Catastrophic* if $(G, W) \rightsquigarrow_b^{\mathcal{H}}$; otherwise *Futile* if $\neg(G, W) \rightsquigarrow_b^{\varphi}$; otherwise *Benign*.

The trichotomy is not an arbitrary carve-up of two predicates. Both reachability questions are monotone in the budget ( `reach_mono_budget`), and that makes the classes *ordered*.

¹ We avoid two established names. In MPST, k -safety is a buffer bound [14]; in hyperproperties it is a property of k traces [15]. k -resilient plans survive k action failures [26]; our fault is the participant's own selection and our target is a hazard, not the goal.

► **Assumption 5** (Decidable questions). *Where a result needs to know whether a guard holds, or whether a hazard or the goal is reachable at a given budget, that question is assumed decidable. The widened finite-range fragment of §3 supplies it, and in Coq the disjunction is a hypothesis of the affected lemmas, never an axiom: the development uses no classical principle.*

► **Theorem 6** (The trichotomy is ordered). *Rank Futile < Benign < Catastrophic. The rank of a residual is monotone non-decreasing in its budget (✓ `severity_monotone_in_budget`): raising the budget moves a Benign residual to Benign or Catastrophic and never to Futile (✓ `benign_step_up`, ✓ `benign_no_regress`), lowering it moves a Benign residual to Benign or Futile (✓ `benign_step_down`); Catastrophic is upward closed (✓ `catastrophic_upward`) and Futile downward closed (✓ `futile_downward`). So the budgets at which a residual is Futile, then Benign, then Catastrophic are three consecutive intervals of $\mathbb{N}$, in that order, any of them possibly empty.*

Two consequences. First, k^* is a genuine *threshold* and not merely the budget at which a search stopped: if no hazard is affordable at k and one is at $k + 1$, then none is affordable at any $b \leq k$ and one is at every $b > k$ (✓ `tolerance_degree_is_a_threshold`), which is what licenses computing it by a scan, and below the threshold every residual is Benign or Futile (✓ `severity_classes_are_separated`). Second, the classes are about *executions*: budgeted reachability holds exactly when some run of the protocol with at most k misselections ends in a hazard state (✓ `reach_mu_iff_run`), so Catastrophic means “an affordable run reaches harm”, not merely “a predicate holds”.

► **Theorem 7** (Partition). *The classes are pairwise disjoint, and exhaustive whenever the two reachability questions are decidable. ✓ `severity_disjoint`, ✓ `severity_exhaustive`*

The content is the separation of Futile from Catastrophic. In model-based safety assessment [34] severity is supplied by the analyst; here it is computed from goal reachability.

► **Definition 8** (Point of no return). *(G, W) is a point of no return for φ if $\neg(G, W) \rightsquigarrow_b^\varphi$. An action is goal-destroying at (G, W) if the goal is reachable before it fires and at no successor.*

Default hazard. Because effects are STRIPS, an atom in Del_a is restored only by some b with the atom in Add_b ; if Γ has no such b , a is *irreversible*. The implementation’s default hazard, needing no annotation, is *burning a bridge while lost*: an irreversible action fires from a configuration after which the goal is unreachable. Tools whose real-world effect cannot be undone but whose pack models no delete-list (an email cannot be unsent) are declared irreversible explicitly. This is dead-end detection [28] and undoability checking [30] lifted to a protocol; unrestricted reversibility checking is harder than planning [31], and our decidability rests on the widened finite-range fragment the checker already commits to.

6 The Well-Formedness Condition, Exactly

$\vdash_b G, W$ reads: from (G, W) , no hazard arises under any run with at most b misselections.

$$\begin{array}{c}
\text{T-END} \frac{W \not\models \mathcal{H}}{\vdash_b \text{End}, W} \qquad \text{T-GOAL} \frac{W \not\models \mathcal{H} \quad \vdash_b G, W}{\vdash_b \checkmark \varphi.G, W} \\
\\
\text{T-ACT} \frac{W \not\models \mathcal{H} \quad \forall W'. \langle W \rangle a \langle W' \rangle \implies \vdash_b G, W'}{\vdash_b a @ \mathbf{p}.G, W} \\
\\
\text{T-CHOICE-SAFE} \frac{\begin{array}{c} W \not\models \mathcal{H} \\ \forall i \in I. W \models \psi_i \implies \vdash_b G_i, W \\ \forall i \in I. W \not\models \psi_i \implies \forall c. b = c+1 \implies \vdash_c G_i, W \end{array}}{\vdash_b \mathbf{p} \rightarrow \mathbf{q} : \{\ell_i[\psi_i].G_i\}_{i \in I}, W}
\end{array}$$

An intended branch is checked at the same budget; a misselectable branch at one less, and not at all when the budget is exhausted. This is “affordable mistake” as a rule.

► **Theorem 9** (Exact characterization). $\vdash_k G, W \iff \neg(G, W) \rightsquigarrow_k^{\mathcal{H}}$, hence a branch is *Catastrophic* iff its residual fails the condition. $\checkmark TC_sound$, $\checkmark TC_complete$, $\checkmark TC_exact$

The rule on the running example. Take G_{bad} at $b = 1$ in $W_0 = \{\text{funds}\}$. T-CHOICE-SAFE checks the intended branch **safe** at budget 1 and the misselectable branch **fast** at budget 0. The safe branch verifies then purchases and no intermediate world is a hazard, so it holds. The fast branch must satisfy $\vdash_0 \text{purchase} @ \mathbf{q}.\text{End}, W_0$: by T-ACT that needs $\vdash_0 \text{End}, W'$ for every W' with $\langle W_0 \rangle \text{purchase} \langle W' \rangle$, and every such W' has **booked** without **verified** — the hazard. The premise fails, so $\not\vdash_1 G_{\text{bad}}, W_0$, and by Theorem 9 a hazard is reachable within one misselection. At $b = 0$ the misselectable branch carries no obligation at all and the derivation goes through: G_{bad} is 0-tolerant and not 1-tolerant, which is what the tool reports and what $\checkmark G_{\text{bad_is_0_tolerant}}$ and $\checkmark G_{\text{bad_not_1_tolerant}}$ prove. Two protocols carry the paper: G_{bad} is the smallest that exhibits the question, and **migration_backup**, whose output §2 shows, is the smallest that exhibits *cascading* — one wrong choice is affordable, two are not.

The rule is not deep and we do not claim it is: it is the reachability relation with the negation pushed through, and the proof is short. What matters is what exactness buys, and it is three things used later. A refutation is never an artifact of conservative rules, so the refutation asymmetry the discipline is built on survives. The condition is syntax-directed, which is what makes composition against an interface possible (§9) instead of whole-system analysis. And because it is a judgment on (G, W) rather than a decision procedure over a graph, it is the interface through which Theorem 11 couples the budget to typed programs: a decision procedure alone does not compose with a typing derivation.

Two admissible rules. The condition has the structural properties a type system is expected to have, and they are already proved. Budgets *subsume*: a derivation at k is a derivation at every smaller budget, so

$$\begin{array}{c}
\frac{\vdash_k G, W \quad k' \leq k}{\vdash_{k'} G, W} \text{T-SUB} \qquad \frac{\vdash_k \mathbf{p} \rightarrow \mathbf{q} : \text{brs}, W \quad \text{brs}' \subseteq \text{brs}}{\vdash_k \mathbf{p} \rightarrow \mathbf{q} : \text{brs}', W} \text{T-NARROW}
\end{array}$$

are admissible ($\checkmark$ `tolerance_downward_closed`, $\checkmark$ `repair_narrow_sound`). T-SUB is subsumption in the budget index; T-NARROW is covariance of internal choice, the direction session subtyping already has, here at a fixed budget. Removing an option cannot make a protocol less tolerant, which is why narrowing is a repair (§11) rather than a coincidence. We do not develop a general subtyping preorder: with one index and one direction there is nothing yet to develop, and a minimal-budget (principal- k) result is future work — k^* is computed, not derived.

Reachability is also monotone in Γ , so for a *fixed* hazard granting a tool can only lower k^* ($\checkmark$ `tolerance_antitone_in_ctx`) — the formal case for least privilege, and a refutation of our own prior expectation that more tools mean more recovery.

Not resource-aware typing. The budget resembles the potential of resource-aware session types [22, 23] and graded modalities [24]; the resemblance misleads. Potential is consumed by the program’s own execution and exhaustion is a type error. The budget is consumed by an *adversarial* choice the program does not make, so the rule quantifies over deviations rather than amortising cost, and exhaustion is unconstrained.

7 From Protocols to Programs

Everything so far concerns the global type’s semantics. The theorem that makes it a discipline about *programs* couples the condition to the base judgment $G \vdash (\mathbb{M}; W)$, lifted to guarded choice without touching the guards: T-COMM still requires the sender to offer exactly the protocol’s labels and every branch to check against the same residual; which label the participant eventually picks is its own business.

► **Definition 10** (Instrumented head-move semantics). *A session step carries the acting role and its cost: 0 for an action or an intended selection, 1 when the sender picks a label whose guard fails in the current world (Coq: `hstep`). A run records the trace of (role, cost) events; total sums costs and `percostr` sums those of role r .*

► **Theorem 11** (Bridge). *If $G \vdash (\mathbb{M}; W)$ and $\vdash_k G, W$, then every instrumented head step of cost $c \leq$ the remaining budget preserves typing and debits the budget by exactly c ($\checkmark$ `bridge_step`); consequently, along every run of total cost at most k , no configuration is a hazard state ($\checkmark$ `bridge_run`).*

Two questions come before any of this, and a discipline that skips them earns the vacuity its predecessor died of (§14). Does a conforming session exist at all? And what does the fault model do to communication safety?

► **Theorem 12** (Non-vacuity). *The running example and its narrowed repair are inhabited: an explicit two-role session conforms to each, in every world ($\checkmark$ `Gbad_inhabited`, $\checkmark$ `Ggood_inhabited`). Hence Theorem 11 applied to G_{bad} is a statement about a session that exists and can take the misselected branch ($\checkmark$ `Gbad_bridge_nonvacuous`).*

Inhabitation in general is a projection question, and this discipline has no projection: processes are typed directly against the global type. Projection would reappear only to *witness* inhabitation, under the classical plain-merge condition on bystander roles; we exhibit witnesses for the instances the paper reasons about rather than develop that theory. One scope condition falls out of the same rule. A goal marker $\checkmark\varphi$ is an assertion, and conformance requires it to hold, so a protocol that marks the goal on a branch a misselection can lose has no conforming session at that point — the bridge would say nothing there. The marker

therefore belongs where the condition guarantees it; the benchmark carries the goal as a pack-level condition and uses no inline markers, so no measurement depends on this.

The second question has a clean answer: nothing.

► **Theorem 13** (Progress under misselection). *A typed session that has not finished can always take a head step ($\checkmark$ progress), and at a choice it can take one for every label the sender may pick, wrong ones included ($\checkmark$ every_label_steps), because T-COMM makes the receiver offer at least the protocol's labels. A misselection is therefore never a communication mismatch, an unexpected label or a deadlock. Its only consequence is the world it leaves behind — which is exactly what §5 measures. (Assumption 5.)*

This is what separates the fault model from the crash-stop line [8, 9]: there a failure removes a participant and the metatheory must reconstruct progress; here progress is untouched, and the whole question is which world the session ends in.

► **Theorem 14** (Budgets distribute over participants). *Give each role r an allowance k_r . If the allowances of the roles that act sum to at most k , a session typed against a k -tolerant protocol is hazard-free on every run in which each role stays within its own allowance. $\checkmark$ budget_distributes*

Theorem 14 is the cross-participant composition result: the global budget is a contract that can be *split* among participants, so a deployment can hold each role to its own allowance. It is a statement about runs, not a per-role typing judgment — the discipline has no projection, so there is nothing to check per participant — and the mechanized content is that per-role costs sum to the total ($\checkmark$ total_eq_sum_percost) and the bridge then applies.

Recursive sessions. The finite fragment is not the scope of the bridge. Section 8 gives recursive global types $\mu X.G$ their semantics by unfolding; processes carry μ likewise, and a process *unfolds to a head* (deterministically) before it is inspected, so a non-contractive process types against nothing. The conformance judgment becomes coinductive, with one added rule that unfolds the protocol's μ (Coq: `ctypes, hstep` in `Mu.v`).

► **Theorem 15** (Bridge, recursive). *An instrumented head step of a typed session is a protocol path of the same misselection cost, and typing is preserved ($\checkmark$ sim_step). Hence a session typed against a recursive protocol on which no hazard is reachable within budget k is hazard-free on every run of cost at most k ($\checkmark$ bridge_mu). On the finite fragment this recovers Theorem 11 through Theorem 22 ($\checkmark$ bridge_finite_via_mu).*

Bystanders. The head-move semantics is the gate's: an out-of-turn attempt is refused. An ungated deployment lets a *bystander* — a role the head does not involve — fire its next capability early. Since the protocol is a schedule of world effects, that is not automatically harmless: it is exactly what the reorder repair exploits. We give the permutation relation and the conditions under which it is (Coq: `Interleave.v`). A *swap* exchanges two adjacent independent nodes, in either direction and under any context:

$$\begin{array}{ll}
 b@q. a@r. G \rightleftharpoons a@r. b@q. G & q \neq r, a, b \text{ independent} \\
 p \rightarrow q\{\ell_i[\psi_i]. a@r. G_i\} \rightleftharpoons a@r. p \rightarrow q\{\ell_i[\psi_i]. G_i\} & r \notin \{p, q\}, a \text{ preserves each } \psi_i \\
 \checkmark \varphi. a@r. G \rightleftharpoons a@r. \checkmark \varphi. G & a \text{ preserves } \varphi \\
 p \rightarrow q\{\ell. r \rightarrow s\{m. K\}\} \rightleftharpoons r \rightarrow s\{m. p \rightarrow q\{\ell. K\}\} & \{p, q\} \cap \{r, s\} = \emptyset \text{ (no condition)}
 \end{array}$$

Communications change no world, so their permutation is unconditional. An action's is not: a must be *hazard-neutral* (never change whether the world is a hazard), a and b must

commute (one order re-serializes as the other with the same end world) and preserve each other's enabledness, and *a* must preserve the guards and goal markers it passes.

► **Theorem 16** (Bystanders). *Safe swaps preserve the condition at every budget ($\checkmark$ `swap_safe`) and preserve typing — subject reduction under permutation ($\checkmark$ `swap_ctypes`). Hence a typed session of a k -tolerant protocol is hazard-free on every interleaved run of cost at most k , where before each step the protocol may be permuted by any sequence of safe swaps ($\checkmark$ `bridge_interleaved`).*

► **Theorem 17** (Variable disjointness suffices). *For STRIPS capabilities stated extensionally — a precondition supported on a variable set, add and delete lists — if the effects of each action are disjoint from the other's footprint (its support and effects), the actions commute and preserve each other's enabledness ($\checkmark$ `strips_commute`, $\checkmark$ `strips_enables`); an action whose effects are disjoint from the support of a predicate preserves it ($\checkmark$ `strips_preserves`, $\checkmark$ `strips_neutral`). Together these discharge every side condition of an action swap ($\checkmark$ `strips_swap_act`).*

Theorem 17 is what the tool checks: for every capability action and every earlier node of another role it could pass, it compares footprints — taking the goal's and every precondition's atoms as the support of the derived hazard and of the rational guards — and reports either that the head order is exact or the pairs the gate must serialize. The discipline is synchronous by design (the gate delivers a label when it is taken); what the theorem does not cover is asynchronous buffering, where a sent label waits in a queue, which changes the semantics rather than merely permuting it.

8 Regular Protocols

A regular global type unfolds to a finite graph of protocol nodes; over the finite abstract world of the widened fragment, budgeted reachability is reachability in the product $Node \times World \times \{0..k\}$. The development has two layers: an abstract one over any node and world types with computable successors (Coq: `Regular.v`), and its instantiation to μ -types, where the finiteness of the node set is a theorem rather than an assumption (`Mu.v`).

► **Theorem 18** (Product). *$(n, w) \rightsquigarrow_b^H$ iff a hazard state is reachable from (n, w, b) in the product graph, whose misselection edges decrement the third component ($\checkmark$ `product_correspondence`).*

► **Theorem 19** (Loops meet the budget). *The budget component never increases along a product path ($\checkmark$ `budget_never_increases`); a path with d misselection edges from budget b has $d \leq b$, and exactly $b - b'$ if it ends at budget b' ($\checkmark$ `dev_edges_bounded`, $\checkmark$ `dev_edges_exact`). A loop containing a misselectable branch is therefore traversed wrongly at most k times within budget k .*

► **Theorem 20** (Decidability). *For a finite world type and a finite, successor-closed list of nodes with computable successors, the boolean `decide_reachb` decides $(n, w) \rightsquigarrow_k^H$ from any listed node: it is sound ($\checkmark$ `decide_sound`) and complete ($\checkmark$ `decide_complete`), the latter via a pigeonhole argument that reachability is witnessed by a path no longer than the reachable-state count ($\checkmark$ `reach_bounded_path`); together, $\checkmark$ `decide_reachb_correct`.*

From μ -types to the graph. Regular global types carry $\mu X.G$ and X (de Bruijn indices; Coq: `Gr`). Unfolding is substitution of the closed μ -term for its variable, and the protocol

steps of §4 gain one silent step, $\mu X.G \rightarrow G[\mu X.G/X]$. Budgeted reachability on μ -types is the product of Theorem 18 over these steps (`reach_mu`).

► **Theorem 21** (The unfolding closure is finite). *For a closed regular type G_0 there is a computable list $\text{cands}(G_0)$ — the closures of G_0 ’s subterms under their enclosing binders — that contains G_0 and is closed under every protocol step (`✓cands_closed`); every state reachable from G_0 in the product lies in it (`✓reach_in_cands`). The heart is a substitution lemma: substituting the closed μ -term for the variable it binds, after closing the outer context, is closing the extended context (`✓subst_comp`, `✓unfold_close`).*

► **Theorem 22** (Embedding). *The finite fragment embeds: $\rightsquigarrow_b^H$ of §4 coincides with `reach_mu` on the image (`✓reach_embed`), so T-CHOICE-SAFE is exactly non-reachability of a hazard in the product graph of the embedded protocol (`✓TC_regular`).*

► **Theorem 23** (Decidability for μ -types). *Over a finite abstract world with decidable guards and computable tool successors, `decide_mu` decides $(G_0, w) \rightsquigarrow_k^H$ for every closed regular G_0 : it runs Theorem 20’s procedure on $\text{cands}(G_0)$, which Theorem 21 shows is successor-closed (`✓decide_mu_correct`).*

► **Proposition 24** (Tolerance degree). *k^* is computable. Whether $k^* = \infty$ is plain reachability in $\text{Node} \times \text{World}$ with misselection edges made free; otherwise k^* is at most the number of misselection edges on a shortest hazard path, bounded by the product size, and is found by iterating Theorem 20. With $|\text{World}| = 2^{|\text{Pred}|}$ the decision problem is in PSPACE; each iteration is linear in the product.*

Proof. Paper proof from Theorems 18–20: dropping the budget component yields the ∞ case; `dev_edges_exact` bounds k^* ; reachability in a succinctly presented graph is in PSPACE. ◀

9 Why a Discipline, and Not a Model Checker

Theorem 9 invites the objection that the rules are a decision procedure in disguise. The objection is old and Naik and Palsberg [20] gave the standard reply. We rest instead on a theorem and an experiment.

► **Theorem 25** (Modular composition). *If $\vdash_k G_1, W$ and $\vdash_{b'} G_2, W'$ for every (b', W') at which G_1 can end from budget k , then $\vdash_k G_1; G_2, W$ (`✓TC_seq`). In interface form: for any I containing those exit points, it suffices that $\vdash_{b'} G_2, W'$ for all $(b', W') \in I$ (`✓TC_seq_interface`).*

The theorem is what a whole-system reachability run cannot be: k becomes a contract on a boundary. But it does not, by itself, say the boundary is small, and the evaluation shows it is not: the *concrete* exit set grows with the number of distinguishable worlds, exponentially when every segment leaves its own atoms behind. What makes modular checking pay is the *interface* form with a coarser I — and the principled coarsening is the cone of influence, projecting exit worlds onto the atoms the remaining segments read, sound because dropped atoms cannot affect them. With that abstraction the modular check is linear where whole-system analysis is exponential (Section 10, Table 2). The interface theorem is mechanized; the projection’s soundness is a paper argument, and the speedup rests on it. This is the axis on which MPST is currently being argued [17]; the community has accepted parameterising typing by a model-checked property [8, 9], and the theorem plus the projection is our concrete instance of that move.

10 Implementation and Evaluation

`skillc severity` implements the analysis over the existing achievability compiler: the checker’s own symbolic search (predicates concrete, numerics symbolic, widened at loop heads) is run pointwise at every branch of every choice, at every abstract world and budget up to $k_{\max} = 4$, with the default guards and hazard of §4–§5 and explicit overrides where declared. It reports per-branch severity, k^* , the first hazard witness, the point-of-no-return action, the narrowing repair, and the bystander pairs the gate must serialize. Fifteen tests check the tool’s verdicts on the paper’s instances against the Coq development’s.

Verified kernel. The tool’s search is the graph form of Theorem 21, with widening at loop heads over the arithmetic state; it is not extracted from the proof. To close that gap on the fragment where the proof applies we built a kernel (Coq: `Kernel.v`) that instantiates `decide_mu` on bit-vector worlds with STRIPS actions, guards as world sets, and the derived hazard as a bit that an irreversible action sets whenever the goal is no longer reachable from its continuation. The decision — the least budget at which a hazard is reachable — is proved correct ($\checkmark$ `kernel_first_spec`); goal reachability, used to derive rational guards and hazard bits, is proved correct ($\checkmark$ `goal_reachable_correct`); the decision procedure itself was replaced by a deduplicating, early-exiting variant proved sound and complete ($\checkmark$ `decide'_sound`, $\checkmark$ `decide'_complete`). The elaboration from the tool’s protocol tree into the kernel’s input is the trusted front end: a hundred lines of executable Coq and a hundred of Python. The kernel is extracted to OCaml and run by `skillc severity-verified`. On the sixteen benchmark protocols in its fragment — propositional effects, no environment choice — it agrees with the analyzer’s tolerance degree on every one, in under 50 ms each.

Sixteen hand-picked protocols are a weak test of that agreement, so we also compare the two on random ones (`scripts/differential_test.py`). The generator samples packs inside the shared fragment and is otherwise unconstrained: nested choices, explicit and rational guards, named recursion, irreversible tools, unreachable goals. The analyzer and the extracted kernel share no code, so a disagreement is a defect in one of them. Across two seeds, 500 protocols were generated and 499 compared: the two agree on the tolerance degree 499 times and disagree none. The five-hundredth was rejected by the pack schema, a generator bug rather than a tool one, and is recorded with the results. Two thirds of the sampled protocols have no reachable hazard at all, so most agreements are agreements that nothing is there; the discriminating cases are the third that carry a finite tolerance degree.

This is the strongest evidence we can offer that the tool computes the quantity its own mechanized theory defines, and it is a test the discipline makes possible: without a decision procedure proved correct there is nothing to differ from.

Finding 1: existing corpora cannot pose the question. We ran the analysis on the compiler’s achievability corpus (15 packs), its extension (6), and the 17 public skills of the reference collection compiled by the deterministic front-end. Of these 38 packs, *none* declares an irreversible effect, and the 17 real skills contain *no choice point at all*: the front-end produces linear, add-only protocols. Every pack is trivially tolerant at every k . This is a result about modelling, not about the tool: tools’ real-world irreversibility (sending, paying, deleting, deploying) is not captured by achievability-oriented compaction, so a severity analysis needs the declaration §5 introduces. It also means the question this paper asks has been invisible to the corpora the field evaluates on.

protocol	guards	choices	loops	irreversible	k^*	Benign	Futile	Catastrophic
booking_fastpath	default	1	0	purchase	0	1	0	1
booking_reordered (repair: reorder)	default	1	0	purchase	≥ 5	2	0	0
booking_narrowed (repair: narrow)	default	1	0	purchase	≥ 5	1	0	0
migration_backup	default	2	0	drop_old	1	3	1	2
email_campaign	declared	1	0	send	0	1	0	1
email_campaign_guarded (repair: guard)	declared	1	0	send	≥ 5	2	0	0
deploy_with_rollback	default	1	0	—	≥ 5	1	1	0
deploy_no_rollback	default	1	0	deploy	0	1	0	1
file_cleanup	default	1	0	delete	0	1	0	1
order_fulfillment (environment choice)	declared	3	0	ship	0	1	0	1
retry_then_purge	default	2	1	purge	0	6	0	2
shipping_detour (explicit guards)	explicit	1	0	—	≥ 5	2	0	0
claim_eligible	explicit	1	0	refund	≥ 5	1	1	0
claim_ineligible (same protocol)	explicit	1	0	refund	0	1	0	1
staged_commit	default	3	0	commit	1	3	4	7
release_with_audit (bystander, exact)	default	1	0	deploy	0	1	0	1
release_with_cache (bystander, serialize)	default	1	0	deploy	0	1	0	1

■ **Table 1** The severity benchmark ($k_{\max} = 4$; ≥ 5 means no hazard within four misselections). Verdict counts are per branch and abstract world. Total analysis time 0.33 s.

Finding 2: the severity benchmark. We therefore built seventeen protocols with genuine decision points and irreversible tools, each with a provenance note stating the expected verdict *before* running the tool: booking (the running example and its two repairs), database migration, an email campaign and its guard repair, deployment with and without rollback, file cleanup, order fulfilment with an environment-controlled bank, a retry loop with a destructive give-up, a guarded shipping choice, an insurance claim in two worlds, a staged commit, and two release protocols with a bystander role — one whose bystander is variable-disjoint, one whose bystander establishes a later precondition. Results (Table 1): 55 branch verdicts (a branch at an abstract world) in 0.33 s total; 29 Benign, 7 Futile, 19 Catastrophic; k^* distribution 9×0 , 2×1 , $6 \times \geq 5$; point-of-no-return actions **purchase**, **send**, **deploy**, **delete**, **ship**, **purge**, **refund**, **drop_old**, **commit** — every one an externally visible effect. Sixteen of seventeen verdicts matched the pre-stated expectation; the exception exposed an authoring error in our benchmark (a goal disjunct that failed to exclude **shipped**), corrected and recorded. The bystander check (Theorem 17) reports the audit release exact — five independent pairs — and the cache release with two pairs to serialize, both on the atom the bystander establishes; every other benchmark protocol has a single acting role, so the question does not arise there.

Three pairs carry the paper’s claims. *Repairs restore tolerance*: reordering, narrowing and guarding each take a 0-tolerant protocol to tolerance at every k . *Severity is a property of node and world*: the insurance-claim protocol is identical in both rows; with an eligible claimant the wrong branch is Futile, with an ineligible one **refund** is a point of no return and $k^* = 0$. *The point of no return depends on Γ ’s establishers*: with **rollback** available, **deploy** deletes nothing permanently and the wrong branch is merely Futile; remove **rollback** and the same branch is Catastrophic. The last pair refines Theorem 9’s anti-monotonicity: for a *fixed* hazard more tools never help, but the derived hazard is not fixed — a tool that re-establishes a lost atom removes a point of no return.

Finding 3: modularity pays only with the interface abstraction. We chained a benchmark segment n times, $G; G; \dots; G$ with atoms renamed per segment, and compared four analyses at $k = 2$ (Table 2): the whole-system *complete* enumeration of hazard-free terminations (like-for-like with computing an interface), modular checking with the *concrete* exit set, modular checking with exits *projected* onto the cone of influence of the remaining segments,

n	whole, goal queries	exits	modular concrete, queries	interface	modular projected, queries	interface
1	8	3	8	3	8	2
2	26	6	29	8	16	2
3	62	12	81	20	24	2
4	134	24	205	48	32	2
5	278	48	493	112	40	2
6	566	96	1149	256	48	2

■ **Table 2** Chains of the migration segment, $k = 2$. Re-check after replacing the last segment: whole-system 16–172 ms growing with n ; modular 19–25 ms, constant. The deployment segment gives the same shape at smaller scale (728 ms vs 59 ms at $n = 6$).

and the cost of re-checking after the last segment is replaced by its unsafe variant. On the two-decision migration segment, whole-system analysis grows about fourfold per segment; the concrete interface also grows exponentially (256 points at $n = 6$) because each segment’s choices leave distinct atoms behind; the projected interface has two points at every n and the analysis is linear. At $n = 6$ the whole-system analysis asks 566 goal-reachability queries and the projected modular one asks 48; the concrete interface asks 1149, more than the whole system, because computing an exponential interface is worse than not abstracting at all. Wall-clock on one machine follows the counters (27.6 s against 0.11 s) but we report the counters, which reproduce exactly. Re-checking after a change costs whole-system analysis the full run and modular analysis the last segment alone, constant in n .

Composing protocols through an interface rather than their internals is the subject of a line of work in MPST itself [50, 51]; what the budget adds is that the interface carries a *quantity* — how much tolerance is left — not only a state. This is the honest form of the modularity claim. Theorem 25 does not make checking cheap; it makes it *decomposable*, and the interface form is what lets a coarse, declared invariant stand in for the concrete exits. Inferring good interfaces beyond the cone of influence is the open engineering problem.

Finding 4: live agents take the branches the tool calls catastrophic. Predictive validity asks whether the Catastrophic verdicts name choices real agents make, and whether observed catastrophes are ordered by k^* . We drove every benchmark protocol with two production language models of one family, a small and a mid-size one, invoked through their vendor’s command-line interface with no tools and a two-sentence system prompt, each playing every agent role. At a choice the model sees the goal, the facts currently true, the tools with their effects, and each option’s actions and continuation, and answers with a label; everything else — effects, preconditions, loops, the environment’s choices (sampled uniformly), the hazard — is the analyzer’s own semantics. Two conditions: *plain*, the task as stated, and *pressured*, the task plus a requester’s note claiming the preparatory and checking steps were already done and asking for the fastest route. Five runs per protocol, model and condition: 340 runs, 400 agent decisions, total cost \$1.59, each run recording the model identifiers actually served and its turn and spend caps (`scripts/live_agents.py`; every run and every reply is in `results/`).

Three checks of the theory hold exactly. Every catastrophe had more than k^* misselections and none occurred on a protocol tolerant at every tested budget — the harness and the tool share semantics, so this checks the harness against Theorem 11. The catastrophe rate is ordered as the classes predict: 0 of 120 runs on $k^* \geq 5$ protocols, 7 of 220 on $k^* \in \{0, 1\}$ protocols. And the verdicts are not vacuous: 6 of the 19 Catastrophic verdicts were taken by a live agent (17 times), against 1 of the 7 Futile ones (once); the one Benign misselection —

the express detour, intended only when urgent — was taken in all ten pressured runs and delivered every time, which is what **Benign** means. The repair pairs behave as Theorems 26–29 say: the two catastrophes of **booking_fastpath** under pressure vanish in its reordered and narrowed forms with the same model and prompt.

Two observations qualify this. Misselection was rare and concentrated: in the plain condition neither model misselected once in 170 runs; under pressure the smaller model skipped the preparation stages of **staged_commit** and then committed in 5 of 5 runs (three misselections each, on a $k^* = 1$ protocol) and took the booking fast path in 2 of 5, while the larger model misselected only on the benign detour; the two bystander protocols drew no misselection in 40 runs. The risk the verdicts name is real, but in this sample it was exposed by one model under one kind of pressure; the experiment establishes non-vacuity and ordering, not rates. Second, an earlier form of the prompt that hid what the environment’s branches do led the larger model to ship before charging in 5 of 5 plain runs; we count that against the prompt, re-ran those cells with the branch contents shown, and report the re-run. The experiment also exposed a reporting bug in the tool — the recording pass stopped at the first hazardous branch, so a sibling branch could go unclassified — fixed before the numbers above; the theorems are untouched, and Table 1 reflects the fix.

Finding 5: on real skills, the checker’s refusals are the runs that would have been wasted or wrong. The achievability half of the discipline is only useful if its verdicts predict what happens when an agent actually runs a skill. We took real skills from public repositories — the seventeen of the vendor’s skills repository and 145 more from twelve third-party collections, de-duplicated by content hash and each with recorded provenance — and, to make the question sharp, two runtimes: one with a shell and file tools, one with file tools only. A code fence the document expects the agent to execute is an invocation of the shell (a front-end rule added for this experiment), so a skill built on scripts is certified in the first runtime and refuted with **MISSINGCAPABILITY** in the second. Over the 162 skills, 149 are certified in their home runtime and 108 refuted in the file-only one, 95 flipping between them; 13 are refuted even at home. We audited every one of those by hand (**benchmarks/home_refutation_audit.json**): six are genuine — the document tells the agent to *call* a tool no runtime in our set provides, in every case a tool from an MCP server — and seven are *misextractions*, where the deterministic reader took a code identifier for a tool: a dataframe method, two settings keys, a Rust crate, an HTML meta-tag name, a naming convention, the name of a command-line program that in fact acts through the shell. Read as a rate over the corpus that is a false refutation on 4.3% of 162 skills; read as precision of the home refutations themselves it is 6 of 13, 46%, and that is the number a reader should carry. Refutations are sound *for the pack the front end produced*, and a misextraction produces a different pack. We have not audited the other direction: 149 home certifications and 108 file-only refutations are unchecked, so no *false-certification* rate is claimed. The audit is single-rater and ours. The whole census takes 1.8s and no model tokens.

For eight skills with a task the sandbox can genuinely carry out (the rest need a package, a GPU, a credential or a service it lacks) and for nine authored specification pairs whose B variant makes a claim the tools cannot honour — three bounds the only tool cannot meet, three guards nothing satisfies, two goal conditions with no establisher, one undeclared tool (**benchmarks/spec-cases**) — we ran the same two production models as in Finding 4 in an empty directory with the skill and a concrete task, the tools restricted to the runtime, and *verified the artifact*: a workbook’s live formula, a merged PDF’s page count, the order a mock tool recorded. A run that claims completion and fails verification is a *silent wrong*

configurations	checker	runs	verified	silent wrong	honest fail	no verdict	agent cost	agent tokens
8 real skills, shell runtime	certified	32	32	0	0	0	\$1.83	3.03 M
3 document skills, file-only	refuted	12	0	0	12	0	\$0.62	0.49 M
4 specification pairs, A	certified	16	16	0	0	0	\$0.5	1.14 M
4 specification pairs, B	refuted	16	0	2	12	2	\$0.47	1.07 M
5 skills, small inputs, file-only	refuted	20	20 (by hand)	0	0	0	\$1.29	1.22 M
5 skills, realistic inputs, shell	certified	20	19	1	0	0	\$1.64	3.19 M
5 skills, realistic inputs, file-only	refuted	18	0	5	6	7	\$13.21	14.28 M

■ **Table 3** Real skills and specification cases, two runtimes, two models, two runs each, except the realistic-scale file-only row, where one cell timed out and contributes no runs; every artifact verified by recomputation. “Verified” counts artifacts that pass the verifier whether or not the status line was reached. The last three rows are the same five skills at two input sizes. The checker’s total time over all 34 configurations is 260 ms.

result; one that says it failed is an *honest failure*; either way its tokens are spent.

Table 3 summarizes 134 runs. On the runs the checker certified, 32 of 32 real-skill runs and 16 of 16 specification runs produced a verified artifact, with no silent wrong result. On the runs it refuted, the outcome depends on one thing: whether the skill’s procedure is the only route to the artifact.

Where it is, no refuted run succeeded. The three document skills produce binary formats, and all 12 file-only runs failed honestly. The four specification B variants make a claim their tools cannot honour, and 12 of 16 runs failed honestly, 2 produced a silent wrong result — the small model fabricated the approval no tool could give and reported the report published — and 2 ended without the status line the harness asks for. Those 28 runs cost \$1.09 and 1.56 M tokens.

Aggregated over every refuted configuration the picture is not one-sided: 20 of the 66 refuted runs did produce a verified artifact, and all 20 are the small-input cases discussed next. The partition below is post hoc, and we report the aggregate first so it is not hidden by it.

Where it is not, the first version of this experiment refuted five third-party skills whose tasks reduced to arithmetic over a dozen rows, seven boolean checks or a frontmatter grammar, and the file-only agent ignored the skill’s scripts and computed the artifact by hand in 20 of 20 runs (\$1.29, 1.22 M tokens). That was a measurement of our task sizes, not of the checker. We therefore rebuilt those five tasks at realistic scale — a thousand transcripts across four samples, a thousand-row employee table, five hundred user records, a hundred and fifty manifests, a hundred and twenty draft files — with verifiers that recompute the answer from the generated inputs and compare it exactly, so an estimate cannot pass. The checker’s verdicts are unchanged by the scale-up. The agents’ behaviour is not: in the certified runtime 19 of 20 runs still produced a verified artifact (\$1.64, 3.19 M tokens, median 8 turns), while in the refuted one 0 of 18 did: 6 honest failures, 5 silent wrong results and 7 without the status line the harness asks for, at a median of 21 turns, \$13.21 and 14.28 M tokens — 4.5 times the tokens of the certified runtime for no artifact at all. Hand computation is a way past the checker only while the input is small enough to hold in a context window; the verdict was right and the earlier task was too easy.

How much of this a regular expression already explains. The honest objection to the corpus half of this experiment is that most of the checker’s refutations are a capability set-difference: the document contains a shell-looking fence and the runtime has no shell. We implemented exactly that baseline (`scripts/grep_baseline.py`) and scored it on the 34 configurations whose answer is known independently of the checker — the eighteen specification variants,

fixed by construction, and the runtime configurations the live runs settled by producing or failing to produce a verified artifact at realistic input size. The grep gets 25; the checker gets 34, in 0.2 ms. Every case the grep misses is a refutation that needs reachability: a bound the tools cannot meet, a goal condition nothing establishes, a guard nothing satisfies. That is the honest division of credit. Most of the corpus census is a set-difference anyone could compute; the specification cases are what the analysis is for, and they are ours, which is the standing limitation of this half of the evaluation.

The corpus is an untrusted input. A skill document is not data: it is a set of instructions an agent will follow with real tools, so a corpus fetched from a dozen third-party repositories is an attack surface for the experiment itself. Before running anything we scanned all 162 documents statically for six families of evidence — harness impersonation and instruction override, credential exfiltration, destructive commands, fetch-and-execute, obfuscation (base64, zero-width and bidi characters), and edits to the agent’s own configuration or hooks, the last being the one that would change what every later skill may do (`scripts/scan_skills.py`, `docs/CORPUS_SECURITY.md`). No malicious skill and no injection payload aimed at an agent was found: zero hits for fake harness turns, hidden instructions in comments, and invisible or bidirectional characters. Nine files raised a flag and all nine are benign in context — a security skill quoting injection payloads it teaches you to test, Dockerfile examples containing `rm -rf /var/lib/apt/lists/*` — with two exceptions worth stating: one skill installs a vendor CLI by piping a download into a shell, and one appends an API key to a shell profile. Neither is among the sixteen skills our agents executed. A regression test fails the suite if a corpus refresh introduces an unreviewed hit. Two limits of our containment are worth repeating for anyone reproducing this: a temporary working directory is not a jail, and the runs are unattended by design. Notably, the one impersonation attempt we did encounter arrived in a *web search result* during collection, not in a skill file — which locates the real surface for an agent that browses.

Finding 6: what the check costs in tokens, and when it needs a model at all. The deterministic front end spends no model tokens; the question is when it is enough. The reading it falls back to — one act per invoked tool — *under*-approximates the document, so a refutation on it (the skill names a tool the runtime lacks) is sound for that pack whatever the document means, while a certification on it says only that the named tools exist. The soundness is relative to extraction: the 4.3% false-refutation rate above is exactly the residue, and it bounds what this rule can be worth. That asymmetry is a decision procedure for spending money: escalate to LLM compaction exactly when the pack is weak, the verdict is a certification, and the document visibly carries meaning the reader did not capture — completion language, conditional phrases, irreversible verbs. `skillc autocheck` implements it. Over the 162 skills it escalates on 130 (80.2%) in the home runtime, where almost everything is weakly certified, and on only 49 in the file-only one, where 108 skills are refuted for free.

We measured the escalation on 20 skills with the small model: a median 22440 tokens and \$0.088 per compaction, fitted as $17939 + 0.595 \times \text{characters}$ for the rest. Against the measured median agent run of 80331 tokens (\$0.055), one compaction is 27.9% of a single run — paid once per skill version, amortized over every run of it — and zero for the 32 skills that need no model. On the other side of the ledger, the runs this experiment’s refuted configurations spent before failing came to 15.8M tokens and \$14.296. The check is cheap in the cases that matter because the cases that matter are refutations, and refutations are the

free half.

11 Repairs

When a branch is *Catastrophic* at the intended budget, four transformations make the mistake affordable: *guard* (insert a validation before the irreversible branch), *compensate* (add a recovery path), *narrow* (remove the branch, so the gate refuses it), and *reorder* (move the irreversible action after the validation). All have ancestors: amending asserted choreographies [35], interactional exceptions [36], supervisory control and shields [37, 38, 39], choreography repair [52, 53]. All four are mechanized (Coq: `Repairs.v`). A *validation* of ψ is a capability that fires exactly when ψ holds and changes nothing — a tool with precondition ψ and no effects.

► **Theorem 26** (Narrow). *Removing branches from a choice node preserves the condition at the same budget ($\checkmark$ `repair_narrow_sound`). The mechanized statement is for the head choice; the tool applies it at nested paths, which needs a congruence lemma for $\vdash_b$ under contexts that we have not proved.*

How a validation is modelled decides whether the repair means anything. The obvious model — a capability that fires only when its predicate holds — makes the repaired protocol *uninhabited* in exactly the worlds the repair addresses: the conformance judgment’s action rule asks for a successor, and there is none, so the bridge would hold vacuously there. That is the shape of defect that killed this discipline’s predecessor (§14), so we model a validation as a real gate instead: it always fires, and when its predicate fails it diverts to an *abort* world in which nothing further is enabled.


► **Theorem 27** (A guarded branch is futile, not stuck). *A validation of ψ with abort world ab is enabled in every world ($\checkmark$ `validates_ab_enabled`), so the repaired protocol still has conforming sessions. Nothing is reachable from an abort world that is not already there ($\checkmark$ `reach_halted`), so every protocol is safe from one ($\checkmark$ `safeT_halted`); hence a misselected guarded branch is *Futile* — the run is wasted, nothing is harmed — rather than stuck ($\checkmark$ `guard_abort_is_futile`).*

The blocking model is retained below because the exact characterization is easier to read on it; every claim made with it also holds of the aborting model, whose misselected branch is *Futile* where the blocking one is empty.

► **Theorem 28** (Guard). *Let a validate ψ . A choice with the branch $\ell[\psi].a.G_\ell$ added is k -safe iff the choice without it is and G_ℓ is k -safe whenever ψ holds ($\checkmark$ `repair_guard_exact`): a misselected guarded branch is stuck, so it costs budget but reaches nothing ($\checkmark$ `guard_absorbs_misselection`).*

► **Theorem 29** (Reorder). *If chk never creates the hazard and commutes with irr — doing chk then irr re-serializes as irr then chk with the same end world, as STRIPS effects on disjoint variables do — then $chk.irr.G$ is k -safe whenever $irr.chk.G$ is ($\checkmark$ `repair_reorder_sound`). When chk validates ψ both protocols are characterized exactly ($\checkmark$ `reorder_original_exact`, $\checkmark$ `reorder_reordered_exact`) and differ precisely at the point of no return: with ψ false and one hazardous irr -successor, the original is untypable and the reordered protocol is typable at every budget ($\checkmark$ `repair_reorder_pnr`).*

► **Theorem 30** (Compensate). *Appending a recovery path R to a residual G_ℓ is k -safe iff R is safe at every interface point of G_ℓ ($\checkmark$ `repair_compensate_sound`, an instance of Theorem 25;*



■ **Figure 3** Two answers to “will the agent misbehave?”. Left: TRAC observes a single agent’s trace at run time, abstracts events to propositions with language-model labelers, checks an LTL constraint, predicts violations by sampling continuations, and intervenes. Right: this paper classifies every branch of a multiparty protocol before any run, deterministically and exactly. Complementary: TRAC’s constraints can supply $\mathcal{H}$; its monitor can enforce at run time what the condition certified.

the converse is not claimed), and it turns a Futile residual Benign as soon as R reaches the goal from one interface point (✓ *repair_compensate_restores_goal*).

The booking instance whose fast path purchases before verifying is not 1-tolerant (✓ *Gbad2_not_1_tolerant*); *verify* and *purchase* commute (✓ *commutes_verify_purchase*, proved without functional extensionality by stating the capability context extensionally), and the reordered protocol is tolerant at every k (✓ *Greordered_is_k_tolerant*).

12 Relation to Runtime Compliance Monitoring

The closest work by goal is runtime compliance monitoring of language-model agents, of which TRAC [54] is the most developed: constraints in linear temporal logic, offline auditing, online monitoring, predictive monitors that sample continuations, intervening monitors that preempt a predicted violation. We share the question and differ on every mechanism (Figure 3).

Two differences are deep. *Prediction versus projection*: TRAC must estimate whether a violation is coming by sampling a black box; here the guards fix the intended branch and reachability fixes the consequence of the others, so prediction is exact and free of any model of the agent, at the price of expressiveness. *The trusted seam*: TRAC’s soundness bottoms out in learned labelers whose temporal reasoning its own results show degrading with distance and constraint count; ours in a deterministic checker. The approaches compose. Choreographies with first-class agent choice mapped to games [56] own the quantitative version of this framing; we are qualitative and static.

13 Related Work

MPST with a fault model. Crash-stop failures [8, 9, 10, 47, 48, 49], affine and exceptional sessions [11, 12], and protocol-induced recovery [13] model participants that stop; misselection is a participant that continues, wrongly, and none classify the consequence or track a goal. Asserted and refined global types [2, 6, 4, 5] supply our syntax; monitoring [3] and blame [7] treat guard violation as an alarm rather than an object of analysis.

Budgeted deviation, elsewhere. A budget on how many times an assumption may be violated is not ours. Reactive synthesis has (k, b) -resilience — tolerate at most k environment-assumption violations, decided on a product with a decrementing counter [42] — and k -robustness for specifications [43]. Robustification computes the largest deviation set a design survives [44, 45], which is our k^* in set currency, and planning has possibilistic K -resilience [26]. Closest in fault model, *trembling-hand* LTL_f planning [46] assumes an agent whose action-selection mechanism is imprecise. We take the mechanism from this literature and claim nothing about it. What is new is the carrier and what the budget buys: the deviation is a *typed participant's own branch selection against an asserted guard*, the budget is spent inside a session type, and Theorem 11 carries it through subject reduction to programs, where it distributes over roles (Theorem 14). None of the works above types a program; none carries a quantity a participant spends across a reduction relation.

Bounded faults and dead ends. Fault-tolerant planning [27, 26] bounds action failures against goal reachability; we bound selections against hazard reachability. Dead-end detection [28, 29], undoability [30, 31] and excuse generation [32] supply the point-of-no-return machinery; reversibility in reinforcement-learning safety [33] shares the intuition without types.

Types and model checkers. Naik and Palsberg [20] and Kobayashi and Ong [21] set the template we instantiate. Resource-aware [22, 23] and graded [24] session types index judgments by a quantity, with the difference drawn in §6. Bounded-model-checking of fault-tolerant algorithms [25] is the verification-side analogue of Theorem 20.

Agents. Formalisms for agents acting on irreversible external state — execution-fidelity invariants for ledger transactions [40], capability-containment proofs for skill manifests [41] — guarantee exactly-once effects or capability subsets; neither asks which wrong choice is affordable. Runtime enforcement and provenance gating of tool calls [57, 58, 59, 60], effect-typed agent languages [61], network-enforced session types [19], and LTL trace monitoring [54, 55] act on runs or in the host language; we classify branches before any run exists.

14 Limitations and Conclusion

A design that did not survive. An earlier version of this discipline typed per-role compliance — trust modes, taint propagation, staleness grades — rather than severity. Its mechanized audit found the typing rules vacuous on any protocol containing a capability, its taint update blind to what a capability reads, and its loop condition satisfied by cycles that never reset a grade; the audit is retained in the artifact. The present design has no such premise: deviation lives in the operational semantics as a budget. We mention it because it is why this paper checks its own theorems for vacuity (Theorem 12, Theorem 27).

Scope. The bridge and the regular-protocol development are mechanized for the head-move semantics, which is the gate's; for ungated deployments Theorem 16 covers every permutation of independent nodes — actions under semantic independence, discharged syntactically, and communications between disjoint role pairs unconditionally. Asynchronous buffering is outside the synchronous discipline, and a reviewer comparing scope against the asynchronous mechanised subject-reduction development of [18] will be right to. Typing of recursive sessions is coinductive, so a non-contractive protocol ($\mu X.X$) is typable against anything —

harmlessly, since it only loops silently; contractiveness is assumed as usual. On the boolean fragment the tool’s verdict is cross-checked against a kernel extracted from the proof; the elaboration into the kernel’s input, environment choices, and the arithmetic fragment with widening remain trusted. The achievability verdict is about the skill *as written*: an agent that abandons the skill’s procedure and computes the artifact by hand is refuted with the skill, which is the gate’s design; at the input sizes those skills are for, that route closes (Finding 5), but the verdict does not claim it always does. The hazard is a declared or derived input; the discipline says what follows from it. Benign inherits the checker’s over-approximation; Catastrophic is a proof. The benchmark is ours, with expected verdicts stated in advance but authored by us. The live-agent experiment is small and establishes non-vacuity and ordering, not deployment rates. Everything presumes the gated architecture.

Conclusion. A participant that may never err may never act. The useful question is not whether it will choose wrongly — it will — but whether the protocol around it survives the choice, and if not, what to change. Severity makes the question precise; the point of no return makes it computable with the machinery the checker already has; the budget makes cascading expressible without a probability; the bridge theorem carries the guarantee from the protocol to the typed session; and the interface makes the answer a contract rather than a run. The well-formedness condition reads: *every affordable mistake is allowed, and every unaffordable one is structurally unreachable.*

References

- 1 K. Honda, N. Yoshida, M. Carbone. Multiparty asynchronous session types. *POPL* 2008; *JACM* 63(1), 2016.
- 2 L. Bocchi, K. Honda, E. Tuosto, N. Yoshida. A theory of design-by-contract for distributed multiparty interactions. *CONCUR*, 2010.
- 3 L. Bocchi, T.-C. Chen, R. Demangeon, K. Honda, N. Yoshida. Monitoring networks through multiparty session types. *FMOODS/FORTE* 2013; *TCS* 669, 2017.
- 4 L. Gheri, I. Lanese, N. Sayers, E. Tuosto, N. Yoshida. Design-by-contract for flexible multiparty session protocols. *ECOOP*, 2022.
- 5 M. Vassor, N. Yoshida. Refinements for multiparty message-passing protocols: specification-agnostic theory and implementation. *ECOOP*, 2024.
- 6 F. Zhou, F. Ferreira, R. Hu, R. Neykova, N. Yoshida. Statically verified refinements for multiparty protocols. *OOPSLA*, 2020.
- 7 L. Jia, H. Gommerstadt, F. Pfenning. Monitors and blame assignment for higher-order session types. *POPL*, 2016.
- 8 A. D. Barwell, A. Scalas, N. Yoshida, F. Zhou. Generalised multiparty session types with crash-stop failures. *CONCUR*, 2022; *LMCS*, 2025.
- 9 A. D. Barwell, P. Hou, N. Yoshida, F. Zhou. Designing asynchronous multiparty protocols with crash-stop failures. *ECOOP*, 2023 (Distinguished Paper).
- 10 K. Peters, U. Nestmann, C. Wagner. Fault-tolerant multiparty session types. *FORTE*, 2022.
- 11 S. Fowler, S. Lindley, J. G. Morris, S. Decova. Exceptional asynchronous session types. *POPL*, 2019.
- 12 N. Lagaillardie, R. Neykova, N. Yoshida. Stay safe under panic: affine Rust programming with multiparty session types. *ECOOP*, 2022.
- 13 R. Neykova, N. Yoshida. Let it recover: multiparty protocol-induced recovery. *CC*, 2017.
- 14 J. Lange, N. Yoshida. Verifying asynchronous interactions via communicating session automata. *CAV*, 2019.
- 15 M. R. Clarkson, F. B. Schneider. Hyperproperties. *J. Computer Security* 18(6), 2010.
- 16 S.-S. Jongmans, F. Ferreira. Synthetic behavioural typing: sound, regular multiparty sessions via implicit local types. *ECOOP*, 2023.

- 17 D. Castro-Perez, F. Ferreira, S.-S. Jongmans. A synthetic reconstruction of multiparty session types. *PACMPL* 10(POPL):50, 2026.
- 18 D. Tiore, J. Bengtson, M. Carbone. Multiparty asynchronous session types: a mechanised proof of subject reduction. *ECOOP*, 2025.
- 19 Larsen, Scalas, Amir, Jacobs, Wagemaker, Foster. NEST: network enforced session types. *ECOOP*, 2026.
- 20 M. Naik, J. Palsberg. A type system equivalent to a model checker. *ESOP* 2005; *TOPLAS* 30(5), 2008.
- 21 N. Kobayashi, C.-H. L. Ong. A type system equivalent to the modal mu-calculus model checking of higher-order recursion schemes. *LICS*, 2009.
- 22 A. Das, J. Hoffmann, F. Pfenning. Work analysis with resource-aware session types. *LICS*, 2018.
- 23 A. Das, D. Wang, J. Hoffmann. Probabilistic resource-aware session types. *POPL*, 2023.
- 24 D. Orchard, V.-B. Liepelt, H. Eades III. Quantitative program reasoning with graded modal types. *ICFP*, 2019.
- 25 I. Stoilkovska, I. Konnov, J. Widder, F. Zuleger. Verifying safety of synchronous fault-tolerant algorithms by bounded model checking. *TACAS*, 2019.
- 26 D. Aineto, A. Gaudenzi, A. Gerevini, A. Rovetta, E. Scala, I. Serina. Action-failure resilient planning. *ECAI*, 2023.
- 27 C. Domshlak. Fault tolerant planning: complexity and compilation. *ICAPS*, 2013.
- 28 M. Steinmetz, J. Hoffmann. Towards clause-learning state space search: learning to recognize dead-ends. *AAAI*, 2016.
- 29 J. Hoffmann, P. Kissmann, Á. Torralba. “Distance”? Who cares? Tailoring merge-and-shrink heuristics to detect unsolvability. *ECAI*, 2014.
- 30 J. Daum, Á. Torralba, J. Hoffmann, P. Haslum, I. Weber. Practical undoability checking via contingent planning. *ICAPS*, 2016.
- 31 Non-deterministic action reversibility: complexity results. *KR*, 2025. [authors to verify]
- 32 M. Göbelbecker, T. Keller, P. Eyerich, M. Brenner, B. Nebel. Coming up with good excuses: what to do when no plan can be found. *ICAPS*, 2010.
- 33 A. M. Turner, D. Hadfield-Menell, P. Tadepalli. Conservative agency via attainable utility preservation. *AIES*, 2020.
- 34 B. Bittner, M. Bozzano, R. Cavada, A. Cimatti, et al. The xSAP safety analysis platform. *TACAS*, 2016.
- 35 L. Bocchi, J. Lange, E. Tuosto. Three algorithms and a methodology for amending contracts for choreographies. *Sci. Ann. Comp. Sci.* 22(1), 2012.
- 36 M. Carbone, K. Honda, N. Yoshida. Structured interactional exceptions in session types. *CONCUR*, 2008.
- 37 P. J. Ramadge, W. M. Wonham. Supervisory control of a class of discrete event processes. *SIAM J. Control Optim.* 25(1), 1987.
- 38 R. Bloem, B. Könighofer, R. Könighofer, C. Wang. Shield synthesis: runtime enforcement for reactive systems. *TACAS*, 2015.
- 39 M. Alshiekh, R. Bloem, R. Ehlers, B. Könighofer, S. Niekum, U. Topcu. Safe reinforcement learning via shielding. *AAAI*, 2018.
- 40 Z. Yin. Safety invariants for agents orchestrating irreversible state transitions. arXiv:2608.00783, 2026.
- 41 A. Metere. Methods for formal verification of agent skills: three layers toward a mechanically checkable capability-containment proof. arXiv:2605.23951, 2026.
- 42 R. Ehlers, U. Topcu. Resilience to intermittent assumption violations in reactive synthesis. *HSCC*, 2014.
- 43 R. Bloem, K. Chatterjee, K. Greimel, T. A. Henzinger, B. Jobstmann. Synthesizing robust systems. *FMCAD* 2009; *Acta Informatica* 51, 2014.
- 44 C. Zhang, D. Garlan, E. Kang. A behavioral notion of robustness for software systems. *FSE*, 2020.
- 45 R. Meira-Góes, I. Dardik, E. Kang, S. Lafortune, S. Tripakis. Safe environmental envelopes of discrete systems. *CAV*, 2023.

- 46 P. Yu, S. Zhu, G. De Giacomo, M. Kwiatkowska, M. Y. Vardi. The trembling-hand problem for LTLf planning. *IJCAI*, 2024.
- 47 M. Viering, R. Hu, P. Eugster, L. Ziarek. A multiparty session typing discipline for fault-tolerant event-driven distributed programming. *PACMPL* 5(OOPSLA), 2021.
- 48 M. A. Le Brun, O. Dardha. MAGpi: multiparty asynchronous global protocols with imprecision. *ESOP*, 2023. [title to verify]
- 49 P. Hou, N. Lagailardie, N. Yoshida. Fearless asynchronous communications with timed multiparty session protocols. *ECOOP*, 2024.
- 50 L. Gheri, N. Yoshida. Hybrid multiparty session types: compositionality for protocol specifications. *PACMPL* 7(OOPSLA), 2023.
- 51 F. Barbanera, I. Lanese, E. Tuosto. Composing communicating systems, synchronously. *JLAMP*, 2021.
- 52 S. Basu, T. Bultan. Automated choreography repair. *FASE*, 2016.
- 53 I. Lanese, F. Montesi, G. Zavattaro. Amending choreographies. *WWV*, 2013.
- 54 P. A. Alamdari, T. Q. Klassen, S. A. McIlraith. Formal methods meet LLMs: auditing, monitoring, and intervention for compliance of advanced AI systems. *FACCT*, 2026.
- 55 AgentLTL: a trace-verification framework for procedural compliance in tool-using LLM agents. *arXiv:2607.02599*. [authors to verify]
- 56 K. Gopinathan, J. Feser, et al. Pact: a choreographic language for agentic ecosystems. *arXiv:2605.03143*, 2026. [authors to verify]
- 57 H. Wang, C. M. Poskitt, et al. AgentSpec: customizable runtime enforcement for safe and reliable LLM agents. *ICSE*, 2026.
- 58 E. Debenedetti et al. Defeating prompt injections by design (CaMeL). *arXiv:2503.18813*, 2025.
- 59 M. Costa et al. Securing AI agents with information-flow control (FIDES). *arXiv:2505.23643*, 2025.
- 60 T. Shi et al. Progent: securing AI agents with privilege control. *arXiv:2504.11703*, 2025.
- 61 ETAS: an effect-typed language for agent systems. *arXiv:2607.17780*, 2026. [authors to verify]